

Assignment 2: Regression Model Building

Andrew G. Dunn¹

¹`andrew.g.dunn@u.northwestern.edu`

Andrew G. Dunn, Northwestern University Predictive Analytics Program

Prepared for PREDICT-410: Regression & Multivariate Analysis.

Formatted using markdown, pandoc, and L^AT_EX. References managed using Bibtex, and pandoc-citeproc.

Simple Linear Regression Models

We'll build upon our EDA, where we found the following continuous variables:

Continuous Variable	Correlation to SalePrice	Prob $> r $ under $H_0: \rho=0$	Number of Observations
GrLivArea	0.70678	<0.0001	2930
GarageArea	0.64040	<0.0001	2929
TotalBsmtSF	0.63228	<0.0001	2929
FirstFlrSF	0.62168	<0.0001	2930
MasVnrArea	0.50828	<0.0001	2907
BsmtFinSF1	0.43291	<0.0001	2929
BsmtUnfSF	0.18286	<0.0001	2929

And first choose to model MasVnrArea, which correlated approximately 0.5 with SalePrice and had almost all of the total observations included. We will use this variable to build a simple linear regression model and comment on the model adequacy.

Second we'll choose GrLivArea, which from our EDA was found to linearly correlate with SalePrice better than MasVnrArea.

Thirdly we will look at the categorical variables;

Categorical Variable	Correlation to SalePrice	Prob $> r $ under $H_0: \rho=0$	Number of Observations
OverallQual	0.79926	<0.0001	2930
GarageCars	0.64788	<0.0001	2929
YearBuilt	0.55843	<0.0001	2930
FullBath	0.54560	<0.0001	2930
YearRemodel	0.53297	<0.0001	2930
GarageYrBlt	0.52697	<0.0001	2771
Fireplaces	0.47456	<0.0001	2930

From this we choose OverallQual, which rates the overall material and finish of the house. It rates on a 10 level Lickert scale [1], 10 being Very Excellent. and 1 being Very Poor.

Model: MasVnrArea predicts SalePrice

The model of interest is

$$\text{SalePrice} = \beta_0 + \beta_1 \text{MasVnrArea} + \epsilon$$

We use the SAS Simple Linear Regression procedure ‘reg’:

```
proc reg;  
  model SalePrice = MasVnrArea;  
run;
```

Which yields the parameters estimates:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	157303	1466.89502	107.24	<0.0001
MasVnrArea	1	226.47763	7.11940	31.81	<0.0001

As such, our fitted model is

$$\text{SalePrice} = 157303 + 226.47763 \times \text{MasVnrArea}$$

Within the context of these variables, the model coefficients indicate that if MasVnrArea was 0 the SalePrice of the house would be \$157,303. We look into our data dictionary to find that MasVnrArea is described ambiguously as ‘Masonry veneer area in square feet’. We state ambiguously because when we examine other variables within the data dictionary we find Exterior 1, Exterior 2, and MasVnrType. These categorical variables show us that there can be multiple exterior finishes, and logically only a subset of those finishes are considered within the MasVnrType variable. At this time, before looking at any graphs, we will suspect to see many observations where MasVnrArea is 0. Therefore, a sale price of \$157,303 is likely reasonable in this model because there are likely many observations in this data set that have a SalePrice and do not have a masonry veneer. This does make us feel quite poorly about the model, as it’s likely that even though we have 2907 observations, only a subset of those observations will be representative of the phenomena we desire to model.

A one unit change in MasVnrArea should be consistent (it is a continuous variable), and the average change in the mean of SalePrice is about \$226.

Both variables have t values that are significantly large ($|t| > 0$) with significant p-values that allow us to reject the null hypothesis and conclude that each of these variables have slope and intercept that are greater than zero.

The model also has some goodness-of-fit information:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	1	4.781879E12	4.781879E12	1011.96	<0.0001
Error	2905	1.372718E13	4725361826		
Corrected Total	2906	1.850905E13			

We look at our F Value and realize that because it is large (greater than 1) the observations and regression will differ from the grand mean. There is some angle, and some linear association with the observations.

From the $\Pr > |t|$ we reject the null hypothesis, and the alternative is that there is a linear relationship between MasVnrArea and SalePrice.

Source	
Root MSE	68741
Dependent Mean	180380
Coeff Var	38.10908
R-Square	0.2584
Adj R-Square	0.2581

We look at our R-Square to see that this regression model only explains $\sim 25\%$ of the variability in SalePrice using MasVnrArea. We will pay attention to the Adjusted R-Square as we continue to build models so that we can compare model performance with consideration to the size of the sample and number of variables are included in the model.

We'll now look at the ODS graphics output for both the Fit Diagnostics and Fit Plot:

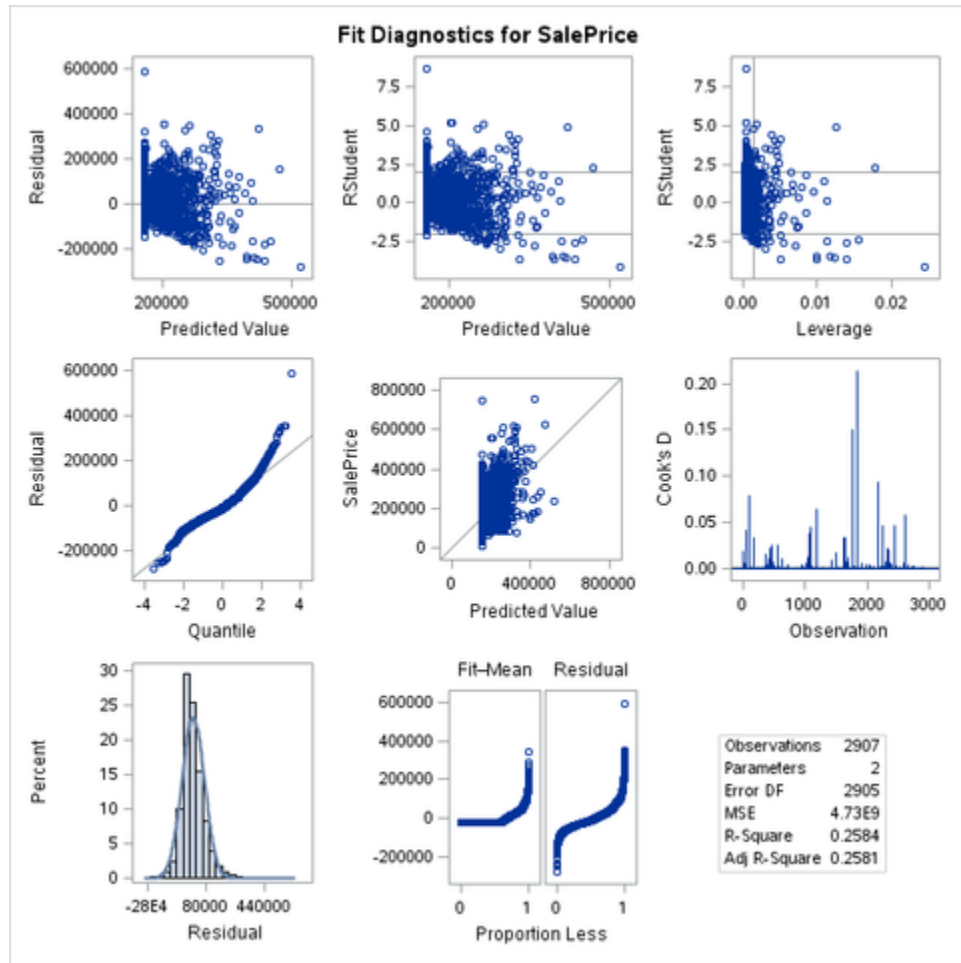


Figure 1: Fit Diagnostics of MasVnrArea vs SalePrice

We assess the normality of residuals by examining the Q-Q plot and observe that the the plot is heavy tailed.

From this we conclude that the observations do not follow the assumed distribution. It is alarming to see that there is a significant outlier in the top right section of the graph. We might be able to live with that outlier, however many observations show up as heavily tailed.

This becomes more of a complication as we observe the histogram of residuals. It appears that the data is mostly normal within this depiction. We at this time have some conflicting feelings based on the two graphs.

We lastly look at the Cook's D. We seek guidance about reading the Cook's D plot and find [2] says: "values of D that are substantially larger than the rest". We also find that there is some conventional wisdom for establishing thresholds for Cook's D using

$$\frac{4}{(N - k - 1)}$$

where N is the number of observations and k is the number of explanatory variables. From this our threshold is computed to be $\frac{4}{2906-1-1} = 0.0013$. We see that there are multiple spikes in the Cook's D over this threshold which should indicate that there are some observations that uniquely influential.

It's safe to say that this model is not homoscedastic due to the above mentioned observations of the graphs. It's important to mention that without the diagnostic graphics we would not have been able to make an assessment of homoscedasticity.




Figure 2: Simple Linear Regression Model of MasVnrArea vs SalePrice

Generally appears to be a positive linear trend, if we ignore that many of the observations for MasVnrArea are at 0. Overall there are some things that we don't like about this model, likely due to many observations of MasVnrArea being zero. We'll look at some other variables to see if we can construct a model that gives us some better feelings about explaining SalePrice.

Model: ‘Best’ Continuous Variable using R-Square Selection

The model of interest is

$$\text{SalePrice} = \beta_0 + \beta_1 \text{'Best'} + \epsilon$$

We use the SAS Simple Linear Regression procedure ‘reg’ with a selection based on the r-square metric:

```
proc reg;  
  model SalePrice = GrLivArea GarageArea TotalBsmtSF FirstFlrSF MasVnrArea BsmtFinSF1 BsmtUnfSF/  
    selection=rsquare start=1 stop=1;  
run;
```

Which yields the model comparison:

R-Square	Variables in Model
0.5006	GrLivArea
0.4085	GarageArea
0.4002	TotalBsmtSF
0.3885	FirstFlrSF
0.2582	MasVnrArea
0.1876	BsmtFinSF1
0.0334	BsmtUnfSF

We then look at the parameter estimates for GrLivArea:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	13290	3269.70277	4.06	<0.0001
GrLivArea	1	111.69400	2.06607	54.06	<0.0001

As such, our fitted model is

$$\text{SalePrice} = 13290 + 111.69400 \times \text{GrLivArea}$$

Within the context of these variables, the model coefficients indicate that if GrLivArea was 0 the SalePrice of the house would be \$13,290. Coming from the variable of MasVnrArea, which from the data dictionary we saw was likely to have many observations be null, we feel stronger about this variable. It is highly likely that there will not be observations in our data set that will be null for this particular variable.

A one unit change in GrLivArea should be consistent (it is a continuous variable), and the average change in the mean of SalePrice is about \$111.70.

Both variables have t values that are significantly large ($|t| > 0$) with significant p-values that allow us to reject the null hypothesis and conclude that each of these variables have slope and intercept that are greater than zero.

The model also has some goodness-of-fit information:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	1	9.3763E12	9.33763E12	2922.59	<0.0001
Error	2928	9.354907E12	3194981962		
Corrected Total	2929	1.869254E13			

We look at our F Value and realize that because it is large (greater than 1) the observations and regression will differ from the grand mean. There is some angle, and some linear association with the observations. From the $\text{Pr} > |t|$ we reject the null hypothesis, and the alternative is that there is a linear relationship between GrLivArea and SalePrice.

Source	
Root MSE	56524
Dependent Mean	180796
Coeff Var	31.26405
R-Square	0.4995
Adj R-Square	0.4994

We look at our R-Square to see that this regression model only explains ~50% of the variability in SalePrice using GrLivArea. We will pay attention to the Adjusted R-Square as we continue to build models so that we can compare model performance with consideration to the size of the sample and number of variables are included in the model.

We'll now look at the ODS graphics output for both the Fit Diagnostics and Fit Plot:

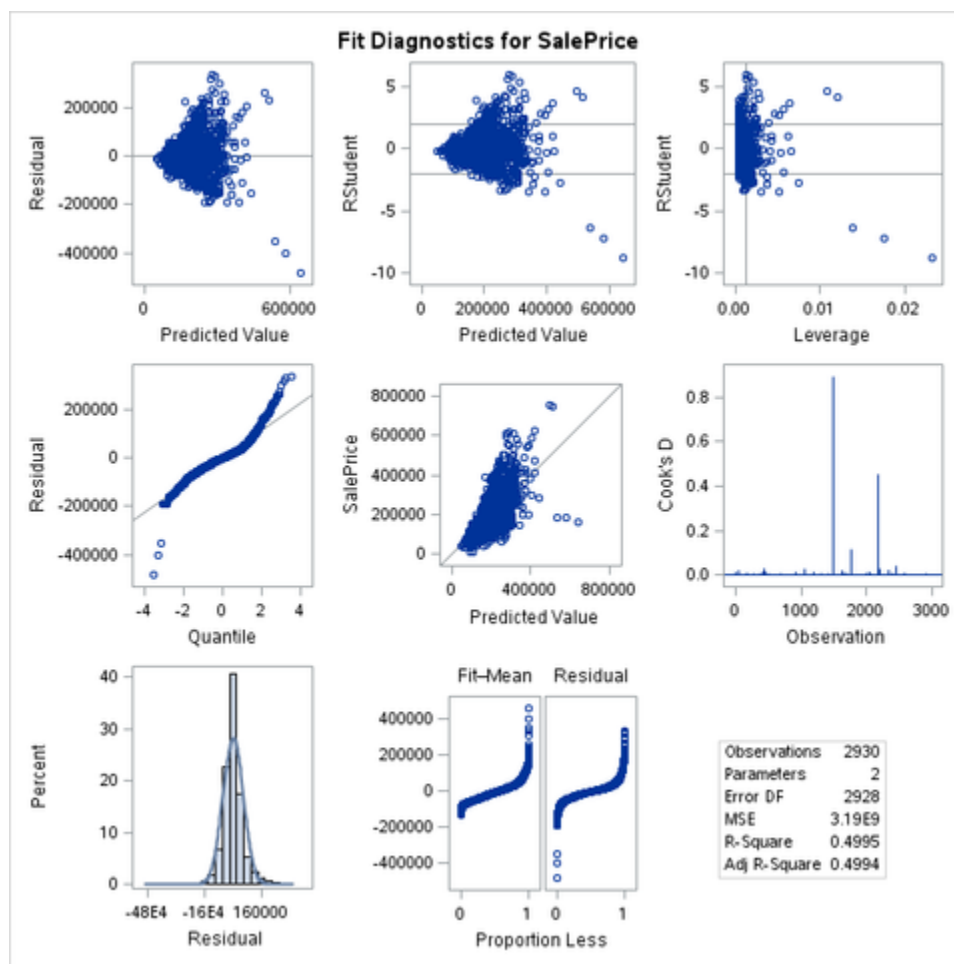


Figure 3: Fit Diagnostics of GrLivArea vs SalePrice

We assess the normality of residuals by examining the Q-Q plot and observe that the plot is heavy tailed. From this we conclude that the observations do not follow the assumed distribution. It is alarming to see that there is a significant outliers in the bottom left of the graph. We might be able to live with some outliers, however many observations show up as heavily tailed.

This becomes more of a complication as we observe the histogram of residuals. It appears that the data is mostly normal within this depiction. We at this time have some conflicting feelings based on the two graphs.

We lastly look at the Cook's D. From the above mentioned equation, our threshold is computed to be $\frac{4}{2930-1-1} = 0.0013$. We see that there are multiple spikes in the Cook's D over this threshold which should indicate that there are some observations that uniquely influential. Particularly three observations, which if we were able to investigate further, which may be the three outliers that are depicted in the bottom left of the Q-Q plot.

It's safe to say that this model is not homoscedastic due to the above mentioned observations of the graphs. It's important to mention that without the diagnostic graphics we would not have been able to make an assessment of homoscedasticity.




Figure 4: Simple Linear Regression Model of GrLivArea vs SalePrice

Generally appears to be a positive linear trend. Compared to the model using MasVnrArea we feel much better about how this model 'looks', even with the above mentioned conclusion that the model isn't homoscedastic. This being the 'best' continuous variable for us to use with a simple linear regression is a bit frightening, even simply examining the Adj R-Square value we can only explain 50% of the variability in our predictor. We may find that a categorical variable, or a more complex model will result in more explanation of variability.

Model: OverallQual predicts SalePrice

The model of interest is

$$\text{SalePrice} = \beta_0 + \beta_1 \text{OverallQual} + \epsilon$$

We use the SAS Simple Linear Regression procedure 'reg':

```
proc reg;  
  model SalePrice = OverallQual;  
run;
```

Which yields the parameters estimates:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	-95004	3933.82223	-24.15	<0.0001
OverallQual	1	45251	628.80511	71.96	<0.0001

As such, our fitted model is

$$\text{SalePrice} = 45251 \times \text{OverallQual} - 95004$$

Within the context of these variables, the model coefficients indicate that if OverallQual was 0 the SalePrice of the house would be \$-95,004. As OverallQual is categorical, and as its a 10 way Lickert scale that begins with 1, it is not reasonable to think of OverallQual being 0.

OverallQual is a categorical variable, a one unit change will result in a much larger jump than the previous continuous variables. If there is a one unit change, our model tells us that the average change in the mean of SalePrice is about \$45,251. Quite a large slope on this due to the categorical variable being between [1, 10].

Both variables have t values that are significantly large ($|t| > 0$) with significant p-values that allow us to reject the null hypothesis and conclude that each of these variables have slope and intercept that are greater than zero.

The model also has some goodness-of-fit information:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	1	1.194116E13	1.194116E13	5178.75	<0.0001
Error	2928	6.751381E12	2305799679		
Corrected Total	2929	1.869254E13			

We look at our F Value and realize that because it is large (greater than 1) the observations and regression will differ from the grand mean. There is some angle, and some linear association with the observations. From the $\text{Pr} > |t|$ we reject the null hypothesis, and the alternative is that there is a linear relationship between OverallQual and SalePrice.

Source	
Root MSE	48019
Dependent Mean	180796
Coeff Var	26.55962
R-Square	0.6388
Adj R-Square	0.6387

We look at our R-Square to see that this regression model only explains ~64% of the variability in SalePrice using OverallQual. We will pay attention to the Adjusted R-Square as we continue to build models so that we can compare model performance with consideration to the size of the sample and number of variables are included in the model.

We'll now look at the ODS graphics output for both the Fit Diagnostics and Fit Plot:

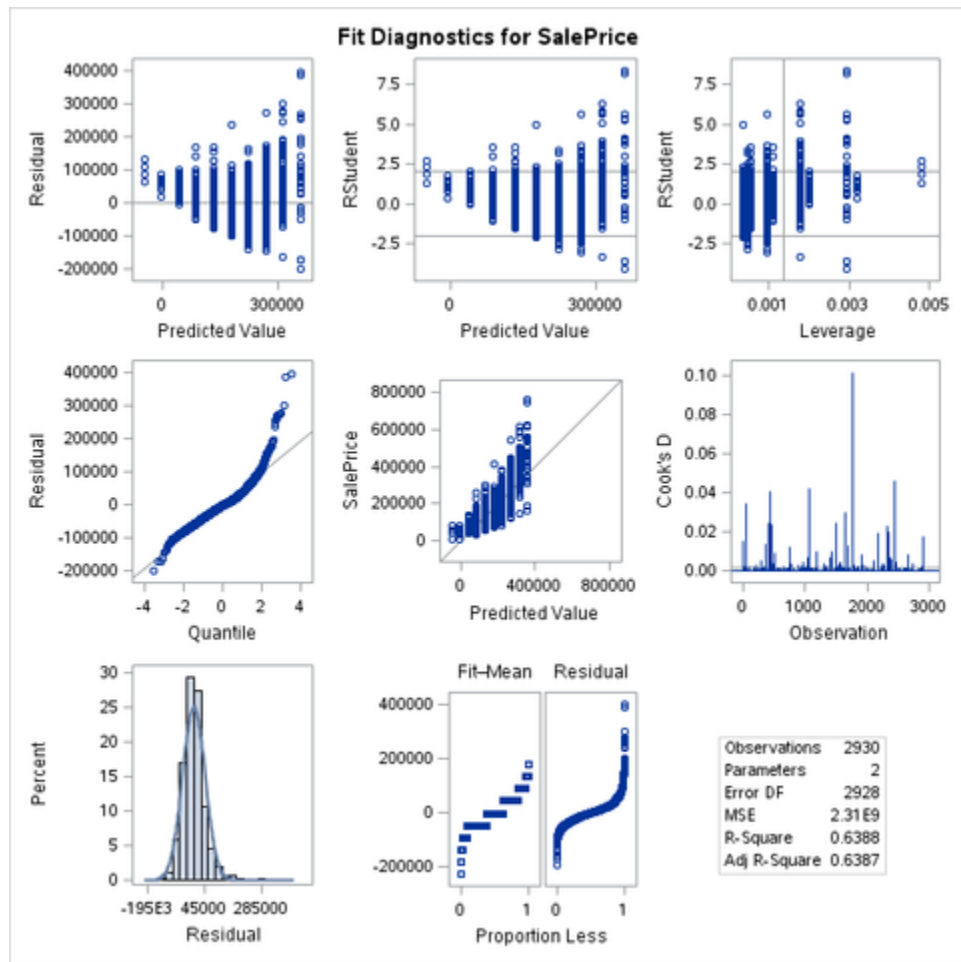


Figure 5: Fit Diagnostics of OverallQual vs SalePrice

We assess the normality of residuals by examining the Q-Q plot and observe that the plot is heavy tailed. From this we conclude that the observations do not follow the assumed distribution. It is alarming to see that there is a significant outliers in the top right section of the graph. We might be able to live with that outlier, however many observations show up as heavily tailed.

This becomes more of a complication as we observe the histogram of residuals. It appears that the data is mostly normal within this depiction. We at this time have some conflicting feelings based on the two graphs.

We lastly look at the Cook's D. From the above mentioned equation, our threshold is computed to be $\frac{4}{2929-1-1} = 0.0013$. We see that there are multiple spikes in the Cook's D over this threshold which should indicate that there are some observations that uniquely influential. In this case there are many observations that are particularly influential.

It's safe to say that this model is not homoscedastic due to the above mentioned observations of the graphs. It's important to mention that without the diagnostic graphics we would not have been able to make an assessment of homoscedasticity.

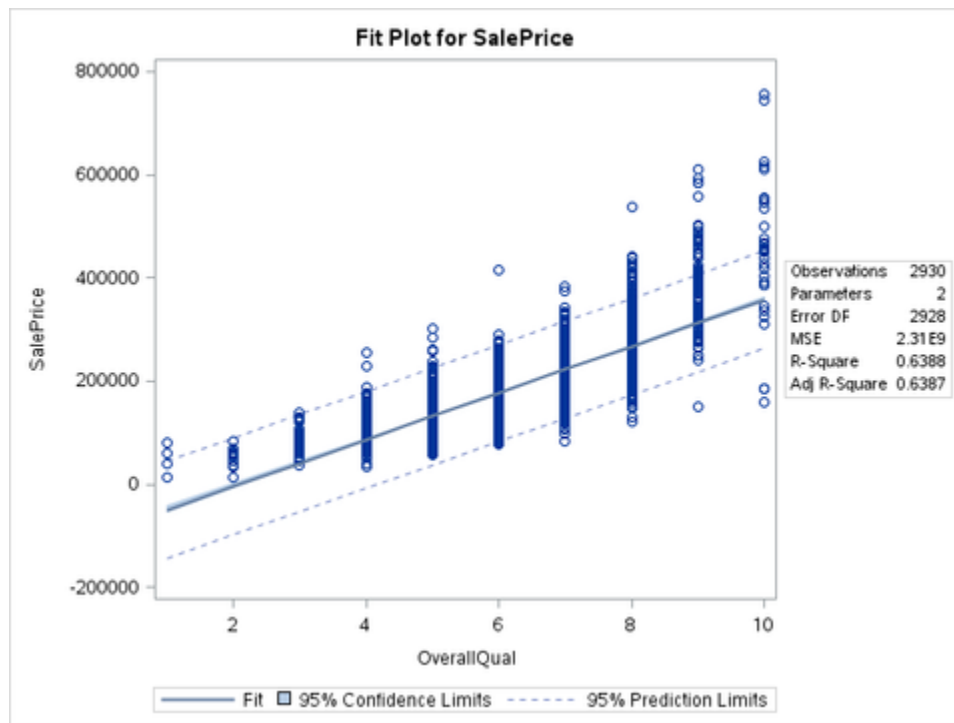


Figure 6: Simple Linear Regression Model of OverallQual vs SalePrice

Generally appears to be a positive linear trend. Compared to the two previous continuous variable models we've built, this model looks quite linear with the largest amount of variability in observation of sale price coming from when a house is rated a 10. Intuitively we'd expect for there to be some high priced outliers that, due to the survey characterization method, would have to be assigned a value of 10. It is interesting to see that there are some outliers at 6, 9, and 10, with some even being low outliers at 10.

If we are to simply compare models based on Adj R-Square values, this model explains the most amount of variability in SalePrice at 64%. Even though this is a categorical variable, and we personally have a bias towards using a continuous variable, we find this to be a highly associative variable.

Simple Linear Model Comparison

We've already commented on how we generally feel about each model towards the end of the section. Here we will compare the diagnostic statistics from each model:

Model	Adj R-Square	F Value
$\text{SalePrice} = 157303 + 226.47763 \times \text{MasVnrArea}$	0.2581	1011.96
$\text{SalePrice} = 13290 + 111.69400 \times \text{GrLivArea}$	0.4994	2922.59
$\text{SalePrice} = 45251 \times \text{OverallQual} - 95004$	0.6387	5178.75

If we solely look at these criteria, then we would conclude that the model that was 'best' based on explanation of variability in SalePrice would be the model that uses OverallQual as the explanatory variable. We can also look at the F Value to see that the model using OverallQual also best fits the population from which the data were sampled.

We've already stated our reservations about using a categorical variable, however within the available variables of this data set, the categorical variable was both linear and highly associative meaning it was a good variable for our modeling approach.

Multiple Linear Regression Models

Model: GrLivArea, MasVnrArea predicts SalePrice

The model of interest is

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \beta_2 \text{MasVnrArea} + \epsilon$$

We use the SAS Simple Linear Regression procedure ‘reg’:

```
proc reg;  
  model SalePrice = GrLivArea MasVnrArea;  
run;
```

Which yields the parameters estimates:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	26547	3141.93715	8.45	<0.0001
GrLivArea	1	94.60302	2.12104	44.60	<0.0001
MasVnrArea	1	118.54695	5.99550	19.77	<0.0001

As such, our fitted model is

$$\text{SalePrice} = 26547 + 94.6032 \times \text{GrLivArea} + 118.54695 \times \text{MasVnrArea}$$

Within the context of these variables, the model coefficients indicate that if GrLivArea was 0, and MasVnrArea was 0 the SalePrice of the house would be \$26,547. We’ve interpreted both GrLivArea and MasVnrArea in the previous sections. It’s unlikely that GrLivArea would ever be 0, however it is possible that MasVnrArea could be 0. The benefit of performing a previous simple linear regression with each of these variables is that we had the opportunity to see some descriptive graphics about the simple model performance. In the event that we didn’t have access to a data dictionary the only way to catch MasVnrArea having 0 in many observations would be to have seen the diagnostic graph.

A one unit change in GrLivArea and MasVnrArea should be consistent (both are continuous variables), and the average change in the mean of SalePrice is about \$213.14.

Both variables have t values that are significantly large ($|t| > 0$) with significant p-values that allow us to reject the null hypothesis and conclude that each of these variables have slope and intercept that are greater than zero.

The model also has some goodness-of-fit information:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	2	1.036256E13	5.181279E12	1846.98	<0.0001
Error	2904	8.146497E12	2805267561		
Corrected Total	2906	1.850905E13			

We look at our F Value and realize that because it is large (greater than 1) the observations and regression

will differ from the grand mean. There is some angle, and some linear association with the observations. From the $\Pr > |t|$ we reject the null hypothesis, and the alternative is that there is a linear relationship between GrLivArea and MasVnrArea, and SalePrice.

Source	
Root MSE	52965
Dependent Mean	180380
Coeff Var	29.36284
R-Square	0.5599
Adj R-Square	0.5596

We look at our R-Square to see that this regression model only explains ~56% of the variability in SalePrice using GrLivArea and MasVnrArea. We will pay attention to the Adjusted R-Square as we continue to build models so that we can compare model performance with consideration to the size of the sample and number of variables are included in the model.

We'll now look at the ODS graphics output for both the Fit Diagnostics and Fit Plot:

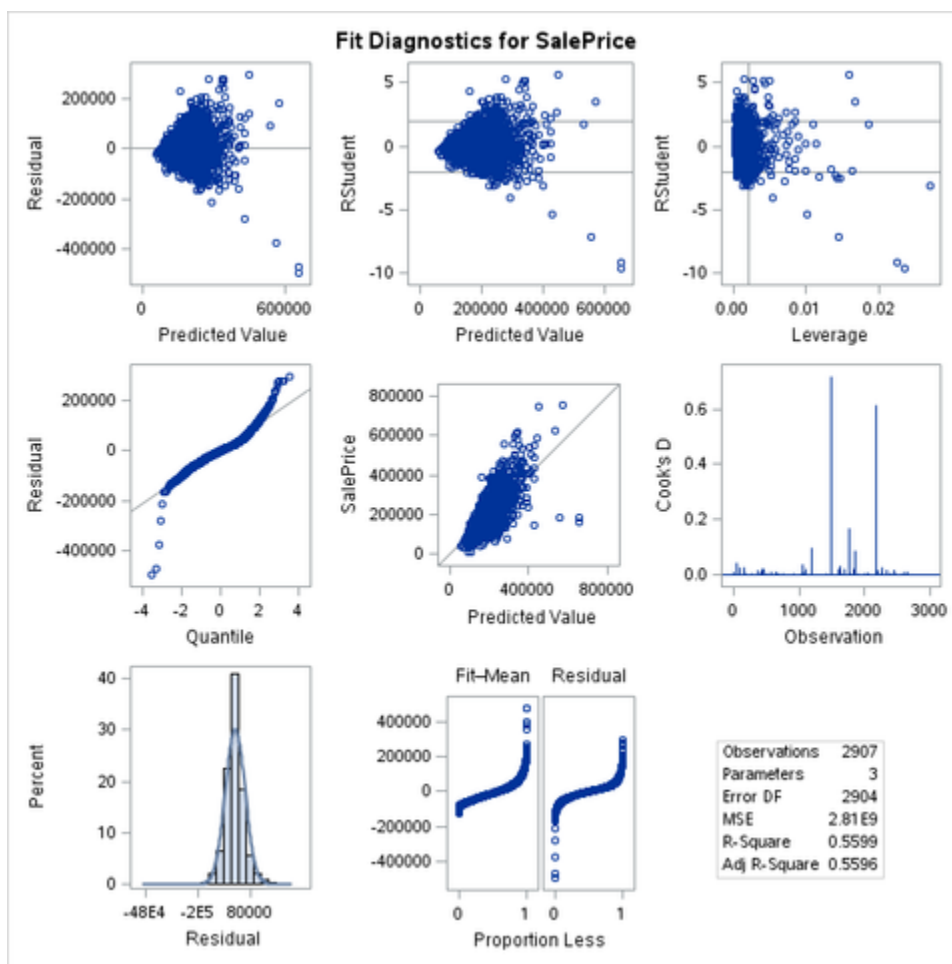


Figure 7: Fit Diagnostics of GrLivArea and MasVnrArea vs SalePrice

We assess the normality of residuals by examining the Q-Q plot and observe that the the plot is heavy tailed. From this we conclude that the observations do not follow the assumed distribution. It is alarming to see that there is significant outliers in the bottom left section of the graph.

This becomes more of a complication as we observe the histogram of residuals. It appears that the data is mostly normal within this depiction. We at this time have some conflicting feelings based on the two graphs.

We lastly look at the Cook's D. From the above mentioned equation, our threshold is computed to be $\frac{4}{2906-1-1} = 0.0013$. We see that there are multiple spikes in the Cook's D over this threshold which should indicate that there are some observations that uniquely influential. In this case there are many observations that are particularly influential.

It's safe to say that this model is not homoscedastic due to the above mentioned observations of the graphs. It's important to mention that without the diagnostic graphics we would not have been able to make an assessment of homoscedasticity.

If we are to simply compare models based on Adj R-Square values, this model explains some amount of variability in SalePrice at 56%. We've already found a simpler model that if we compare just the Adj R-Square values will perform better than this multiple regression model.

Model: GrLivArea + MasVnrArea + ‘Worst’ (BsmtUnfSF) predicts SalePrice

We find that the BsmtUnfSF variable is both continuous and has a very poor correlation to SalePrice at 0.18286.

The model of interest is

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \beta_2 \text{MasVnrArea} + \beta_3 \text{BsmtUnfSF} + \epsilon$$

We use the SAS Simple Linear Regression procedure ‘reg’:

```
proc reg;  
  model SalePrice = GrLivArea MasVnrArea BsmtUnfSF;  
run;
```

Which yields the parameters estimates:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	26547	3141.93715	8.45	<0.0001
GrLivArea	1	94.60302	2.12104	44.60	<0.0001
MasVnrArea	1	118.54695	5.99550	19.77	<0.0001
BsmtUnfSF	1	3.23138	2.30311	1.4	0.1607

As such, our fitted model is

$$\text{SalePrice} = 26547 + 94.6032 \times \text{GrLivArea} + 118.54695 \times \text{MasVnrArea} + 3.23138 \times \text{BsmtUnfSF}$$

Within the context of these variables, the model coefficients indicate that if GrLivArea was 0, MasVnrArea was 0, and BsmtUnfSF was 0 the SalePrice of the house would be \$26,547. We notice that the intercept doesn’t change from the last model we built that excluded the BsmtUnfSF variable.

A one unit change in GrLivArea, MasVnrArea, and BsmtUnfSF should be consistent (all are continuous variables), and the average change in the mean of SalePrice is about \$215.37.

All three variables have t values that are significantly large ($|t| > 0$), both GrLivArea and MasVnrArea have significant p-values where BsmtUnfSF has a p-value that is much higher. We can conclude that the slope for GrLivArea and MasVnrArea is greater than zero, however we have to accept the null hypothesis for BsmtUnfSF, meaning the slope is zero.

The model also has some goodness-of-fit information:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	3	1.035884E13	3.452948E12	1231.02	<0.0001
Error	2902	8.13993E12	2804938144		
Corrected Total	2905	1.849877E13			

We look at our F Value and realize that because it is large (greater than 1) the observations and regression will differ from the grand mean. There is some angle, and some linear association with the observations. From the $\text{Pr} > |t|$ we reject the null hypothesis, and the alternative is that there is a linear relationship

between GrLivArea, MasVnrArea, BsmtUnfSF and SalePrice.

Source	
Root MSE	52962
Dependent Mean	180415
Coeff Var	29.35544
R-Square	0.5600
Adj R-Square	0.5595

We look at our R-Square to see that this regression model only explains ~56% of the variability in SalePrice using GrLivArea, MasVnrArea and BsmtUnfSF. We will pay attention to the Adjusted R-Square as we continue to build models so that we can compare model performance with consideration to the size of the sample and number of variables are included in the model.

We'll now look at the ODS graphics output for both the Fit Diagnostics and Fit Plot:

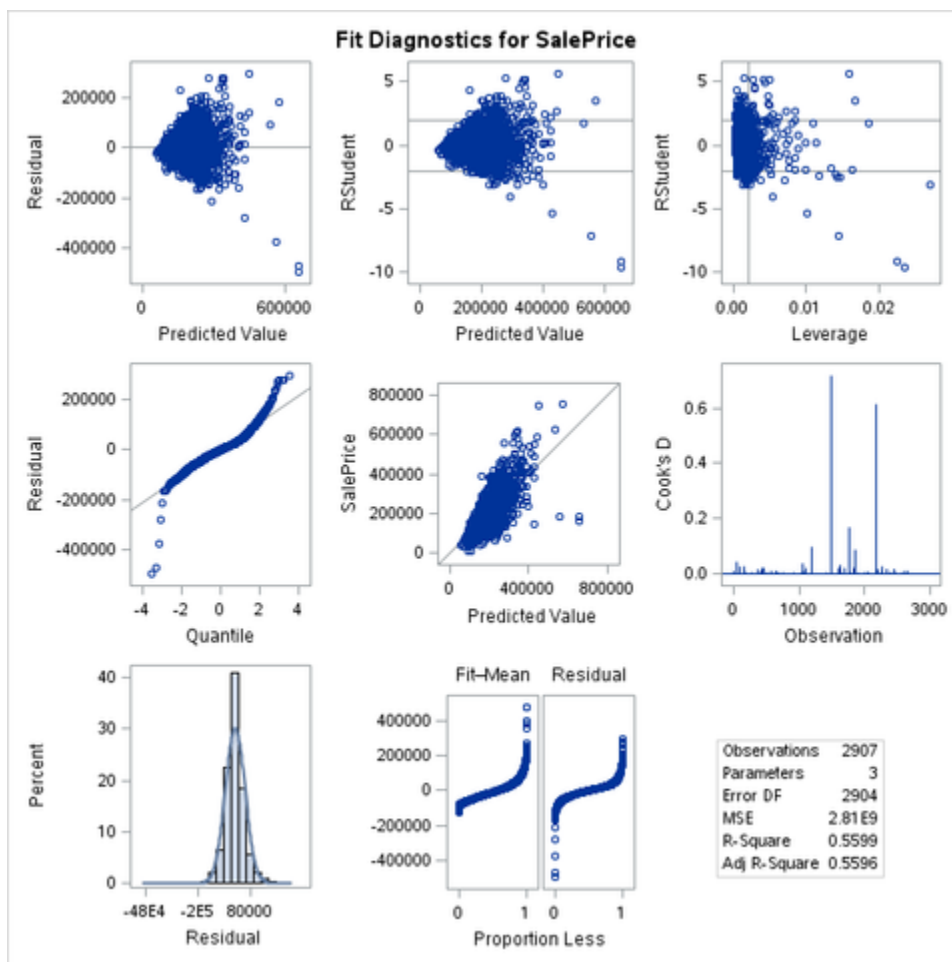


Figure 8: Fit Diagnostics of GrLivArea, MasVnrArea, and BsmtUnfSF vs SalePrice

We assess the normality of residuals by examining the Q-Q plot and observe that the plot is heavy tailed. From this we conclude that the observations do not follow the assumed distribution. It is alarming to see

that there is significant outliers in the bottom left section of the graph.

This becomes more of a complication as we observe the histogram of residuals. It appears that the data is mostly normal within this depiction. We at this time have some conflicting feelings based on the two graphs.

We lastly look at the Cook's D. From the above mentioned equation, our threshold is computed to be $\frac{4}{2905-1-1} = 0.0013$. We see that there are multiple spikes in the Cook's D over this threshold which should indicate that there are some observations that uniquely influential. In this case there are many observations that are particularly influential.

It's safe to say that this model is not homoscedastic due to the above mentioned observations of the graphs. It's important to mention that without the diagnostic graphics we would not have been able to make an assessment of homoscedasticity.

Comparing this model to the last model is quite interesting. We selected an obviously poor explanatory variable and by incorporating it into the model our raw R-Squared score increased. Our Adj R-Square value decreased by a slight amount and our F Value decreased significantly.

It appears that from some perspectives adding more predictor variables, no matter how bad, results in a better R-Square value. This is not necessarily an indication of better 'fit' as much as an indication that if you throw in any garbage to a multivariate model you will explain more of the variability in your predictor variable.

Conclusion / Reflection

We've drawn some comparison between the models that were constructed in this assignment. Primarily we're using the Adj R-Square and F-Value metrics to make those comparisons quantitatively. There is also some qualitative comparisons to be made when constructing a model, for example the Masonry Veneer Square Footage variable would be something to be avoided early on because it only has partial values for the observations in the data set.

This assignment illustrated the importance of visually inspecting model performance, rather than just examining diagnostic statistics.

The next steps in the modeling process for this data set would be to see if there is other data that can be used to begin an attempt at model validation. It may be valuable at this time to present the initial model findings to the business owner. In this case, our initial assessment is that a categorical variable performs the best for explaining variability of Sale Price. If the Business Owner doesn't want us to use a categorical variable, or would rather we use a multi-variable approach to include both continuous and categorical variables it would be best to know that before proceeding to validation.

Procedures

```
title 'Assignment 2';

libname mydata '/scs/crb519/PREDICT_410/SAS_Data/' access=readonly;

* create a temporary variable (data source is read only);
data ames;
    set mydata.ames_housing_data;

* initial examination of the correlation to saleprice;
proc corr data=ames nosimple rank;
    var saleprice;
    with _numeric_;
run;

ods graphics on;

* regression model for SalePrice with predictor MasVnrArea;
proc reg data=ames;
    model SalePrice = MasVnrArea;
run;

* regression model for SalePrice with selection based on rsquare;
proc reg data=ames;
    model SalePrice = GrLivArea GarageArea TotalBsmtSF FirstFlrSF MasVnrArea BsmtFinSF1 BsmtUnfSF /
        selection=rsquare start=1 stop=1;
run;

* regression model for SalePrice with predictor GrLivArea;
proc reg data=ames;
    model SalePrice = GrLivArea;
run;

* regression model for SalePrice with predictor OverallQual;
proc reg data=ames;
    model SalePrice = OverallQual;
run;

* regression model for SalePrice with predictor GrLivArea, MasVnrArea;
proc reg data=ames;
    model SalePrice = GrLivArea MasVnrArea;
run;

* regression model for SalePrice with predictor GrLivArea, MasVnrArea, BsmtUnfSF;
proc reg data=ames;
    model SalePrice = GrLivArea MasVnrArea BsmtUnfSF;
run;

ods graphics off;
```

References

- [1]Wikipedia, “Likert scale — wikipedia, the free encyclopedia.” 2015 [Online]. Available: http://en.wikipedia.org/w/index.php?title=Likert_scale&oldid=653173542
- [2]J. Fox, “Regression diagnostics: An introduction (quantitative applications in the social sciences).” Sage Publications, Inc., Thousand Oaks, CA, 1991.