

Assignment 6: Principal Components in Predictive Modeling

Andrew G. Dunn¹

¹andrew.g.dunn@u.northwestern.edu

Andrew G. Dunn, Northwestern University Predictive Analytics Program

Prepared for PREDICT-410: Regression & Multivariate Analysis.

Formatted using markdown, pandoc, and L^AT_EX. References managed using Bibtex, and pandoc-citeproc.

Goals & Data Examination

Our data set consists of daily closing stock prices for twenty stocks and a large-cap index fund from Vanguard (VV). Our data ranges from January 3rd 2012 to December 31st 2013, a record for each day in our data set (501 days). We don't appear to have any gaps (skipping a day) or null values in our data set. We begin by computing the log return of each stock/index variable. We define return: r_i at a time i , where p_i is the price at time i and $j = (i - 1)$:

$$r_i = \frac{p_i - p_j}{p_j}$$

Where log-return is simply: $\log(r_i)$, and time i is in days. We use return instead of price because it provides a mechanism of normalization, which allows us to have a measurement of all variables in a comparable metric, thus enabling evaluation of analytical relationships amongst two or more variables despite originating from a price series of unequal values [1]. We wish to use the log-returns of the individual stocks to explain the variation in log-returns of the market index (VV). We will explore this concept using both linear regression and principal components analysis as a preconditioner for linear regression analysis.

As with any data set we'll have some basis of contextual knowledge and from that we'll likely be looking to have assumptions validated or broken. For this data set, an example assumption would be that, the context of market sector is a valuable grouping mechanism. Furthermore, we assume to see that sector that are heavily price regulated (banking, energy) would have similar correlations to a large market index.

With this study we plan to build a model where we have the stock values (or a normalized re-expression) as independent variables, and the large market index Vanguard (VV) as the dependent variable.

Correlation between log-return and Market Index

When we perform the PearsonCorr we get back a wide dataset, meaning that the variables in this vector are in a left to right orientation. This isn't the best for examining, and we're going to make some graphics later that will require the data to be transposed into long format. Once transposed we create a separate data set with each ticker and their respective sector. We sort both the transposed to long correlation data set as well as the new ticker and sector data set by ticker, then merge them. We are then able to print a table of the Pearson Correlations of log-return to the Vanguard Index, to include each ticker respective sector:

Observation	Correlation	Ticker	Sector
1	0.63241	AA	Industrial - Metals
2	0.65019	BAC	Banking
3	0.57750	BHI	Oil Field Services
4	0.72090	CVX	Oil Refining
5	0.68952	DD	Industrial - Chemical
6	0.62645	DOW	Industrial - Chemical
7	0.44350	DPS	Soft Drinks
8	0.71216	GS	Banking
9	0.59750	HAL	Oil Field Services
10	0.61080	HES	Oil Refining
11	0.76838	HON	Manufacturing
12	0.58194	HUN	Industrial - Chemical
13	0.65785	JPM	Banking

Observation	Correlation	Ticker	Sector
14	0.59980	KO	Soft Drinks
15	0.76085	MMM	Manufacturing
16	0.47312	MPC	Oil Refining
17	0.50753	PEP	Soft Drinks
18	0.69285	SLB	Oil Field Services
19	0.73357	WFC	Banking
20	0.72111	XOM	Oil Refining

Table 1: Ticker Correlation to Vanguard Index Fund

We want to examine the table above by creating a bar chart where we plot the correlation on the Y-axis, but group the tickers by sector on the X-axis. We will vary the color of each ticker so that it is easier to discern:




Figure 1: Ticker correlation, grouped by sector, to Market Index Vanguard

It is interesting to see that the manufacturing sector has similar correlation, as well as the banking sector. Surprising to us based on first assumption, the oil refining sector doesn't have a similar set of correlations.

We had some initial assumptions about sectors that have heavy price regulation all having similar correlation to a large market identifier. We will now examine the mean correlation by sector, the data table:

Observation	Sector	Type	Frequency	Mean Correlation
1	Banking	1	4	0.68844
2	Industrial - Chemical	1	3	0.63264
3	Industrial - Metals	1	1	0.63241
4	Manufacturing	1	2	0.76461
5	Oil Field Services	1	3	0.62262
6	Oil Refining	1	4	0.63148
7	Soft Drinks	1	3	0.51694

Table 2: Mean Sector Correlation to Vanguard Index Fund

We want to examine the table above by creating a bar chart:




Figure 2: Mean Correlation with Market Index Vanguard by Sector

We observe that the manufacturing sector seems to have the highest correlation, and all other sectors aside

from soft drinks have a similar correlation. Instead of doing the calculation prior, and producing the table above, we could have used special functionality within SAS to perform the mean calculation for us at the time of rendering the graphic.

Principal Components

We create a new data set to exclude the Market Index as when we calculate our principle components we don't want the dependent variable within the data set. We can use a trick when we call the set directive within SAS, passing in the keep statement, as well as 'return_:' using the colon as a wild-card to match all of our previously calculated log-return tickers.

We will use SAS princomp to calculate the eigenvectors as well as produce a scree plot. We'll look at the text output of the princomp procedure, as well as visually examine the scree plot, all with the intent to come up with some criteria for which principle components we plan to take forward into our regression modeling. If we had some arbitrary threshold for variability explanation, we'd be doing some simple calculations with eigenvalues to see how we can exceed that threshold. We'll examine the eigenvalues:

Observation	Eigenvalue	Difference	Proportion	Cumulative
1	9.63645075	8.09792128	0.4818	0.4818
2	1.53852947	0.19109235	0.0769	0.5587
3	1.34743712	0.39975791	0.0674	0.6261
4	0.94767921	0.15217268	0.0474	0.6735
5	0.79550653	0.12909860	0.0398	0.7133
6	0.66640793	0.10798740	0.0333	0.7466
7	0.55842052	0.04567198	0.0279	0.7745
8	0.51274854	0.01590728	0.0256	0.8002
9	0.49684126	0.03250822	0.0248	0.8250
10	0.46433304	0.03089374	0.0232	0.8482
11	0.43343929	0.02568332	0.0217	0.8699
12	0.40775598	0.05667006	0.0204	0.8903
13	0.35108592	0.01597897	0.0176	0.9078
14	0.33510695	0.03813712	0.0168	0.9246
15	0.29696984	0.02068234	0.0148	0.9394
16	0.27628750	0.01692712	0.0138	0.9532
17	0.25936037	0.01730228	0.0130	0.9662
18	0.24205809	0.02020002	0.0121	0.9783
19	0.22185807	0.01013445	0.0111	0.9894
20	0.21172363	-	0.0106	1.0000

Table 3: Eigenvalues of the Correlation Matrix

We now look at a scree plot that is overlay with the cumulative values from the table above:




Figure 3: Scree with Variance Explained (Cumulative)

The Scree plot gives us a very quick and simple mechanism for evaluating the set of calculated principal components and how each one of them contributes to the overall variability within the data set. From both the table and the Scree diagram we can decide on our own strategy for what principal components will be taken to regression modeling. For example, if our threshold for explanation of variability was 80%, we would take the first eight principal components.

Furthermore, we'll perform a sanity check and examine the first two principle component vectors to see if we notice any relationship within the data:




Figure 4: Principal Component 1 and Principal Component 2

Within this graph we notice two clusters. We examine the tickers within each cluster to see if they belong to the same sector and find that not to be true. Initially, we had an expectation that we would see HES, HAL, AA all be within the banking sector. We don't see any way to reason through why these clusters are occurring. With many tickers, it would take a while, we could sit and examine within this graph to see what tickers are directly related (positive on either axis) and what tickers are inversely related (negative on either axis).

Segment PCA Data Set into Train/Test

We'll split the original data set into training and test groups, 70% of the data being selected at random to be within the training group. We will use the merge feature within the SAS data directive to merge the original log-return calculations in with the principle components. On the rows that we flag as part of our training set we assign the response_VV to a new variable train_response.

Regression Model

We fit a regression model using all of the raw predictor variables. We examine the diagnostic output:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	20	0.01790	0.00089510	140.04	< 0.0001
Error	317	0.00203	0.00000639		
Corrected Total	317	0.01993			

Table 4: Analysis of Variance

Source	
Root MSE	0.00253
Dependent Mean	0.00061635
Coeff Var	410.18453
R-Square	0.8983
Adj R-Square	0.8919

Table 5: Model Performance

We see that the R-Square and Adj. R-Square are high. We expected this as we know when we incorporate all variables into the model due to:

- Implicitly, R^2 is a function of the number of parameters in the model. For a nested subset of predictor variables $p_0 < p_1$, i.e. p_1 contains the original p_0 predictor variables and some new predictor variables, R^2 will have a monotonic relationship such that $R^2(p_0) \leq R^2(p_1)$.

We note our F Value so that we can compare with other models, we also note that it is significant and now need to examine the parameter estimates:

Variable	DF	Parameter Estimate	Standard Error	t-Value	Pr > t	Variance Inflation
Intercept	1	0.00008640	0.00014092	0.61	0.5403	0.00000
return_AA	1	0.01769	0.01317	1.34	0.1802	2.11490
return_BAC	1	0.03198	0.01165	2.75	0.0064	3.10927
return_BHI	1	-0.00111	0.01323	-0.08	0.9333	2.62997
return_CVX	1	0.04907	0.02536	1.93	0.0539	3.07524
return_DD	1	0.04674	0.02037	2.29	0.0224	2.51406
return_DOW	1	0.03642	0.01162	3.14	0.0019	1.88893
return_DPS	1	0.03670	0.01679	2.19	0.0295	1.54768
return_GS	1	0.04849	0.01555	3.12	0.0020	3.10450
return_HAL	1	0.00948	0.01466	0.65	0.5184	3.08758

Variable	DF	Parameter Estimate	Standard Error	t-Value	Pr > t	Variance Inflation
return_HES	1	0.00359	0.01092	0.33	0.7425	2.10199
return_HON	1	0.12213	0.01924	6.35	<.0001	2.73505
return_HUN	1	0.02712	0.00836	3.24	0.0013	1.79852
return_JPM	1	0.00902	0.01708	0.53	0.5979	3.36439
return_KO	1	0.07903	0.02226	3.55	0.0004	1.93633
return_MMM	1	0.09796	0.02646	3.70	0.0003	2.98277
return_MPC	1	0.01673	0.00809	2.07	0.0394	1.32999
return_PEP	1	0.02911	0.02231	1.30	0.1929	1.68825
return_SLB	1	0.03776	0.01709	2.21	0.0279	3.13690
return_WFC	1	0.07587	0.01848	4.10	<.0001	2.59492
return_XOM	1	0.05467	0.02697	2.03	0.0435	2.98393

Table 6: Parameter Estimates

It's pretty safe to say that our nose should be detecting the presence of multicollinearity, we see that we have a significant F Value and almost all of the independent variables have a non-significant t-test. We examine the Variance Inflation Factor of the k^{th} predictor as:

$$VIF_k = \frac{1}{1 - R_k^2}$$

where R_k^2 is the R-Square value obtained by regressing the k^{th} predictor on the remaining predictors. VIF_k tells us how much the variance of the estimated regression coefficient β_k is "inflated" by the existence of correlation among the predictor variables in the model. A rule of thumb is if $VIF_k > 10$ then multicollinearity is considered to be high. We see that none of the predictors has a VIF higher than 4, however we see many of the predictors having a close to 4 value (which is considered another lower threshold for VIF calculations). It is generally desired that we achieve a VIF_k value close to 1.0.

Regression Model with Principal Components

We fit a regression model using the eight selected principal components. We examine the diagnostic output:

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	8	0.01776	0.00222	337.13	< 0.0001
Error	329	0.00217	0.00000659		
Corrected Total	337	0.01993			

Table 7: Analysis of Variance

Source	
Root MSE	0.00257
Dependent Mean	0.00061635
Coeff Var	416.36522
R-Square	0.8913
Adj R-Square	0.8886

Table 8: Model Performance

Initially we note that the F Value has increased over our last model, albiet the R-Square and Adj. R-Square have decreased slightly. We examine the parameter estimates:

Variable	DF	Parameter Estimate	Standard Error	t-Value	Pr > t	Variance Inflation
Intercept	1	0.00075978	0.00014045	5.41	<.0001	0.00000
Prin1	1	0.00231	0.00004519	51.05	<.0001	1.00527
Prin2	1	0.00032245	0.00011425	2.82	0.0051	1.00868
Prin3	1	0.00070635	0.00012322	5.73	<.0001	1.00861
Prin4	1	0.00030481	0.00014536	2.10	0.0368	1.00636
Prin5	1	-0.00017356	0.00015516	-1.12	0.2641	1.00297
Prin6	1	0.00000315	0.00017108	0.02	0.9853	1.00766
Prin7	1	-0.00010331	0.00018604	-0.56	0.5791	1.02315
Prin8	1	-0.00040760	0.00020293	-2.01	0.0454	1.02271

Table 9: Parameter Estimates

We observe that the first three independent variables are significant. This is an improvement over the last model. We observe that independent variable four through eight are not significant, but in examining the VIF_k of each predictor variable we observe them all to be close to one.

Model Comparison & Conclusions

We observe that between the two models, the model based on the log-return variables has significant signs of multicollinearity. The model based on the first eight principal components has less of an indication of multicollinearity. The principal components model has a larger F Value, while both models have a significant measure of F Value. We'll compute the Mean Square Error (MSE) and Mean Absolute Error (MAE) to further compare the models. We calculate the MSE as:

$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$$

Where \hat{Y} is a vector of n predictions, and Y is the vector of the true values. MSE assesses the quality of an estimator or set of predictions in terms of its variation and degree of bias. We calculate the MAE as:

$$MAE = \frac{1}{n} \sum_{i=1}^n |\hat{Y}_i - Y_i|$$

Where \hat{Y} is a vector of n predictions, and Y is the vector of the true values. MAE is an average of the absolute error within the model. In both cases we want to see that the calculated values are minimized. We will perform both calculations on the training and test data:

Observation	Training	MSE_1	MAE_1	MSE_2	MAE_2
1	0	0.000009306	0.002144904	0.000009677	0.002179249
2	1	0.000005994	0.001902032	0.000006410	0.001975239

Table 10: Mean Square Error & Mean Absolute Error

Overall the difference between models when looking at MSE and MAE is small. The percent difference, to normalize the comparison, is:

MSE training	3.9867%
MSE Test	6.9403%
MAE Training	1.6012%
MAE Test	3.8489%

Overall, we have to favor the second model for a variety of reasons:

- Significantly fewer parameters, making the model simpler to interpret. Although the caveat of for interpretation is explaining dimensionality reduction to a business owner.
- Less obvious presence of multicollinearity with low VIF_k on each parameter
- Better Goodness-of-Fit with a higher F-Value, and relatively the same R-Square, and Adj. R-Square

With the first model having obvious multicollinearity using PCA as a preprocessor for regression gave us the ability to address multicollinearity by reducing the amount of parameters we use in the model (while still accounting for a tunable (our choice) amount of variability from the data set.

Procedures

```
* William T. Mickelson ;  
* 5-5-2015 ;  
* portfolio_pca.sas;
```

```
libname mydata "/scs/wtm926/" access=readonly;
```

```
data temp;  
    set mydata.stock_portfolio_data;  
run;
```

```
proc print data=temp(obs=10); run; quit;  
proc sort data=temp; by date; run; quit;
```

```
data temp;  
    set temp;
```

```
* Compute the log-returns - log of the ratio of today's price to yesterday's price;  
* Note that the data needs to be sorted in the correct  
    direction in order for us to compute the correct return;  
return_AA = log(AA/lag1(AA));  
return_BAC = log(BAC/lag1(BAC));  
return_BHI = log(BHI/lag1(BHI));  
return_CVX = log(CVX/lag1(CVX));  
return_DD = log(DD/lag1(DD));  
return_DOW = log(DOW/lag1(DOW));  
return_DPS = log(DPS/lag1(DPS));  
return_GS = log(GS/lag1(GS));  
return_HAL = log(HAL/lag1(HAL));  
return_HES = log(HES/lag1(HES));  
return_HON = log(HON/lag1(HON));  
return_HUN = log(HUN/lag1(HUN));  
return_JPM = log(JPM/lag1(JPM));  
return_KO = log(KO/lag1(KO));  
return_MMM = log(MMM/lag1(MMM));  
return_MPC = log(MPC/lag1(MPC));  
return_PEP = log(PEP/lag1(PEP));  
return_SLB = log(SLB/lag1(SLB));  
return_WFC = log(WFC/lag1(WFC));  
return_XOM = log(XOM/lag1(XOM));  
*return_VV = log(VV/lag1(VV));  
response_VV = log(VV/lag1(VV));  
run;
```

```
proc print data=temp(obs=10); run; quit;
```

```
* We can use ODS TRACE to print out all of the data sets available to ODS for a particular SAS procedure  
* We can also look these data sets up in the SAS User's Guide in the chapter for the selected procedure  
*ods trace on;  
ods output PearsonCorr=portfolio_correlations;  
proc corr data=temp;  
*var return: with response_VV;  
var return_;;
```

```

with response_VV;
run; quit;
*ods trace off;

proc print data=portfolio_correlations; run; quit;

data wide_correlations;
  set portfolio_correlations (keep=return_);
run;

* Note that wide_correlations is a 'wide' data set and we need a 'long' data set;
* SAS has two 'standard' data formats - wide and long;
* We can use PROC TRANSPOSE to convert data from one format to the other;
proc transpose data=wide_correlations out=long_correlations;
run; quit;

data long_correlations;
  set long_correlations;
  tkr = substr(_NAME_,8,3);
  drop _NAME_;
  rename COL1=correlation;
run;

proc print data=long_correlations; run; quit;

* Merge on sector id and make a colored bar plot;
data sector;
input tkr $ 1-3 sector $ 4-35;
datalines;
AA Industrial - Metals
BAC Banking
BHI Oil Field Services
CVX Oil Refining
DD Industrial - Chemical
DOW Industrial - Chemical
DPS Soft Drinks
GS Banking
HAL Oil Field Services
HES Oil Refining
HON Manufacturing
HUN Industrial - Chemical
JPM Banking
KO Soft Drinks
MMM Manufacturing
MPC Oil Refining
PEP Soft Drinks
SLB Oil Field Services
WFC Banking
XOM Oil Refining
VV Market Index

```

```

;
run;

proc print data=sector; run; quit;

proc sort data=sector; by tkr; run;
proc sort data=long_correlations; by tkr; run;

data long_correlations;
  merge long_correlations (in=a) sector (in=b);
  by tkr;
  if (a=1) and (b=1);
run;

proc print data=long_correlations; run; quit;

* Make Grouped Bar Plot;
* p. 48 Statistical Graphics Procedures By Example;
ods graphics on;
title 'Correlations with the Market Index';
proc sgplot data=long_correlations;
  format correlation 3.2;
  vbar tkr / response=correlation group=sector groupdisplay=cluster datalabel;
run; quit;
ods graphics off;

* Still not the correct graphic. We want the tickers grouped and color coded by sector;
* We want ticker labels directly under the x-axis and sector labels under the ticker
  labels denoting each group. Looks like we have an open SAS graphics project.;
ods graphics on;
title 'Correlations with the Market Index';
proc sgplot data=long_correlations;
  format correlation 3.2;
  vbar sector / response=correlation group=tkr groupdisplay=cluster datalabel;
run; quit;
ods graphics off;

* SAS can produce bar plots by sector of the mean correlation;
proc means data=long_correlations nway noprint;
  class sector;
  var correlation;
  output out=mean_correlation mean(correlation)=mean_correlation;
run; quit;

proc print data=mean_correlation; run;

ods graphics on;
title 'Mean Correlations with the Market Index by Sector';

```

```

proc sgplot data=mean_correlation;
  format mean_correlation 3.2;
  vbar sector / response=mean_correlation datalabel;
run; quit;
ods graphics off;

* Note that we have been using PROC SGPLOT to display a data summary, and hence we have not
been able to make the display that we want. In reality PROC SGPLOT is designed to take an
input data set, perform some routine data summaries, and display that output. Unfortunately,
routine data summaries are typically frequency counts for discrete data or averages for
continuous data. Here is an example of the default use of PROC SGPLOT.;
ods graphics on;
title 'Mean Correlations with the Market Index by Sector - SGPLOT COMPUTES MEANS';
proc sgplot data=long_correlations;
  format correlation 3.2;
  vbar sector / response=correlation stat=mean datalabel;
run; quit;
ods graphics off;

* Reset title statement to blank;
title '';

*****;
* Begin Modeling;
*****;
* Note that we do not want the response variable in the data used to compute the
principal components;

data return_data;
  set temp (keep= return_);
  * What happens when I put this keep statement in the set statement?;
  * Look it up in The Little SAS Book;
run;

proc print data=return_data(obs=10); run;

*****;
* Compute Principal Components;
*****;
ods graphics on;
proc princomp data=return_data out=pca_output outstat=eigenvalues plots=scree(unpackpanel);
run; quit;
ods graphics off;
* Notice that PROC PRINCOMP produces a lot of output;
* How many principal components should we keep?;
* Do the principal components have any interpretability?;
* Can we display that interpretability using graphics?;

proc print data=pca_output(obs=10); run;

proc print data=eigenvalues(where=( _TYPE_='SCORE' )); run;
* Display the two plots and the Eigenvalue table from the output;

```

```

* Plot the first two eigenvectors;
data pca2;
  set eigenvectors(where=(_NAME_ in ('Prin1','Prin2')));
  drop _TYPE_ ;
run;

proc print data=pca2; run;

proc transpose data=pca2 out=long_pca; run; quit;
proc print data=long_pca; run;

data long_pca;
  set long_pca;
  format tkr $3.;
  tkr = substr(_NAME_,8,3);
  drop _NAME_;
run;

proc print data=long_pca; run;

* Plot the first two principal components;
ods graphics on;
proc sgplot data=long_pca;
scatter x=Prin1 y=Prin2 / datalabel=tkr;
run; quit;
ods graphics off;

*****;
* Create a training data set and a testing data set from the PCA output ;
* Note that we will use a SAS shortcut to keep both of these 'datasets' in one ;
* data set that we will call cv_data (cross-validation data). ;
*****;
data cv_data;
  merge pca_output temp(keep=response_VV);
  * No BY statement needed here. We are going to append a column in its current order;
  * generate a uniform(0,1) random variable with seed set to 123;
  u = uniform(123);
  if (u < 0.70) then train = 1;
  else train = 0;

  if (train=1) then train_response=response_VV;
  else train_response=.;

run;

proc print data=cv_data(obs=10); run;

* You can double check this merge by printing out the original data;
proc print data=temp(keep=response_VV obs=10); run; quit;

```



```

*****;
* Fit a regression model using the raw predictor variables;
*****;
* Fit a regression model using all of the raw predictor variables and VV as the response variable;

ods graphics on;
proc reg data=cv_data;
model train_response = return_ / vif ;
output out=model1_output predicted=Yhat ;
run; quit;
ods graphics off;
* Examine the Goodness-Of-Fit for this model. How well does it fit? Are there any problems?;

*****;
* Fit a regression model using the rotated predictor variables (Principal Component Scores) ;
*****;
* Now fit a regression model using your selected number of principal components and VV as
  the response variable;
* Examine the Goodness-Of-Fit for this model. How well does it fit? Are there any problems?;

ods graphics on;
proc reg data=cv_data;
model train_response = prin1-prin8 / vif ;
output out=model2_output predicted=Yhat ;
run; quit;
ods graphics off;

*****;
* Compare fit and predictive accuracy of the two fitted models ;
*****;
* Use the Mean Square Error (MSE) and the Mean Absolute Error (MAE) metrics for your comparison.;

proc print data=model1_output(obs=10); run;

* Model 1;
data model1_output;
  set model1_output;
  square_error = (response_VV - Yhat)**2;
  absolute_error = abs(response_VV - Yhat);
run;

proc means data=model1_output nway noprint;
class train;
var square_error absolute_error;
output out=model1_error
  mean(square_error)=MSE_1
  mean(absolute_error)=MAE_1;
run; quit;

```

```

proc print data=model1_error; run;

* Model 2;
data model2_output;
  set model2_output;
  square_error = (response_VV - Yhat)**2;
  absolute_error = abs(response_VV - Yhat);
run;

proc means data=model2_output nway noprint;
  class train;
  var square_error absolute_error;
  output out=model2_error
    mean(square_error)=MSE_2
    mean(absolute_error)=MAE_2;
run; quit;

proc print data=model2_error; run;

* Merge them together in one table;
data error_table;
  merge model1_error(drop= _TYPE_ _FREQ_) model2_error(drop= _TYPE_ _FREQ_);
  by train;
run;

proc print data=error_table; run;
* Which model fits better? Did we benefit from using PCA? ;
* In our analysis did we do anything that might not be 100% kosher? ;
* Model 2 has higher accuracy both in-sample and out-of-sample, hence we prefer Model 2;

```

References

[1]Quantitivity, “Why log returns.” 2011 [Online]. Available: <https://quantitivity.wordpress.com/2011/02/21/why-log-returns/>