Study

My study notes will draw heavily from the required texts and multimedia. I will also draw from external sources that I find to be adept at explaining a particular topic. If something is referenced here, it is because I found it to be very useful in understanding a topic.

Standard Mathematical and Statistical Notation

Notes below are from the following sources; [Bhatti 2011a].

Vector and Matrix Notation

- A *scalar* is a number. *Scalars* are represented by lower case letters from the beginning of the alphabet such as a, b, c etc.
- A *vector* is a $n \times 1$ array defined with the mathematical operations of addition and multiplication. The standard convention is for all vectors to be column vectors, i.e. they are 'long' with n rows and 1 column. Vectors are represented as **bold** faced lower case letters frequently from the end of the alphabet, such as $\mathbf{x}, \mathbf{u}$, and $\mathbf{v}$. the ith entry of a vector $\mathbf{u}$ is denoted by $\mathbf{u}[i] = u_i$.
- A *matrix* is a $n \times m$ array defined with the mathematical operations of addition and multiplication. Matricies are represented by a bold face upper cased letter such as $\mathbf{A}, \mathbf{W}, \mathbf{X}$, etc. The (i, j)th entry of a matrix $\mathbf{A}$ is denoted by $\mathbf{A}[i, j] = a_{ij}$.
- The transpose of a $(n \times 1)$ column vector $\mathbf{a}$ is the $(1 \times n)$ row vector $\mathbf{a}^T = [a_1 \dots a_n]$. Sometimes the transpose $\mathbf{a}^T$ is denoted by $\mathbf{a}'$.
- The transpose of a $(n \times m)$ matrix $\mathbf{A}$ is the $(m \times n)$ matrix $\mathbf{A}^T$ where $\mathbf{A}[i, j] = \mathbf{A}^T[j, i]$. When a matrix is transposed, the rows become the columns and the columns become the rows.
- It is preferred to use the T notation $\mathbf{a}^T$ instead of the "prime notation" $\mathbf{a}'$.

Random Variable Notation

Random variables are how you develop calculus based probability theorey.Random variables are the unkown statistical experiement that generate data. Statistical theory is based upon the concept of random sampling.

- Random variables are denoted by capital letters from the end of the alphabet such as U, V, X, Y, or Z.
- The *observed value* of a random variable is denoted by the lower cased counterpart such as u, v, x, y, or z.
- When we have a *random sample* of independent and identically distributed (iid) random variables, we will index the variables in a set such as $X_1, X_2, \dots, X_n$ for the random variables and $x_1, x_2, \dots, x_n$ for the observed values.
- Random variables are used to devleop statistical estimators. Observes values of random variables are used to compute statistical estimates.
- Random variable notation can become convoluded when we move to multivariate random variables. Pay attention to how an author presents these concepts in text.

'Distribution'

The term *distribution* is used throughout all statistical applications and discussions. Loosly speaking, the term distribution is meant to describe how a group of values are related to either each other or to the range of values on which they are defined (their *support*).

The term *distribution* is used rather sloppily. If you don't understand the context you wont understand the use. The term *distribution* is mapped to many related concepts. In general the term *distribution* is related to the characterization of a random variable, or data generated by a random variable.

There are many mathematical notations for characterizing a statistical distribution. The choice of characterization will depend on the context and the existence of the characterization. A random variable can be characterized by any of the following functions.

- The *cumulative distribution function* (cdf), denoted by $F(x) = Pr(X \leq x)$. the cdf will exist for all random variables, and in general is why we use the term "distribution" so looslely throughout statistics. cdf exists for all random variables. From data you can always estimate a distribution function.
- The *probability density function* (pdf) for continuous random variables, denoted by $f(x)$, or the *probability mass function* (pmf) for discrete random variables, denoted by $p(x)$. Note that neither of these functions are guaranteed to exist. A random variable that can be described with a cdf will not always possess a pdf or pmf.
- Transformation functions such as the moment generating function $m(t) = \mathbb{E}[\exp(tX)]$ and the characteristic function $\phi(t) = \mathbb{E}[\exp(itX)]$. Transformation functions are not used when working with data, they may be used to develop conceptual underpinnings of modeling.
- Specialized representations for particular applications such as the *hazard fucntion* $h(t) = \frac{f(t)}{S(t)}$ and the *survival function* $S(t) = 1 - F(t)$ used in Survival Analysis. Survival function is a simple map of the cdf. The hazard function allows you to get a generic representation of a survival function.
- In data analysis distributions can be analyzed using the empirical cdf, the histogram, the Quantile-Quantile plot, and the Kolmogorov-Smirnov test.
- If you need to assess the distribution of residuals in linear regression and compare that to the assumption that they are normally distributed.

Mathematical Expectation

Mathematical Expertation is the theoretical averaging of a random variable with respect to its distribution function. In this sense the pdf of pmf act as a weight function that allows you to find the "center" of the distribution.

For a continuous random variable X with pdf function $f(x)$, the mathematical expectation of X can be computed by

$$\mathbb{E}[X] = \int x f(x) dx$$

For a discrete random variable X with pmf function $p(x) = Pr(X = x)$, the mathematical expectation of X can be computed by

$$\mathbb{E}[X] = \sum_x x p(x)$$

$\mathbb{E}[X]$ is also referred to as the first moment of X.

Expectation, Variance, and Covariance as Mathematical Operators

Let X denote a random variable. Consider the affine transformation $aX + b$.

- $\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$
- $\text{Var}[aX + b] = a^2\text{Var}[X]$

Let X and Y be random variables with a joint distribution function. (In the continuous case we would denote this joint distribution function by the join density function $f(x, y)$.) Consider the linera transformtions aX and bY.

- $\mathbb{E}[aX + bY] = a\mathbb{E}[X] + b\mathbb{E}[Y]$
- $\text{Var}[aX + bY] = a^2\text{Var}[X] + b^2\text{Var}[Y] + ab\text{Cov}[X, Y]$

Here the reader should note that in general $\text{Cov}[aX + b, cY + d] = ac\text{Cov}[X, Y]$. If X and Y are independent random variables, then $\text{Cov}[X, Y] = 0$. The converse of this statement is not true except when both X and Y are normally distributed. In general $\text{Cov}[X, Y] = 0$ does not imply that X and Y are indepdented random variables.

Statistical Assumptions for Ordinary Least Squares Regression

Notes below are from the following sources; [Bhatti 2011b].

- In Ordinary Lease Squares (OLS) regression we wish to model a continuous random variable Y (the response variable) given a set of *predictor variables* $X_1, X_2, \ldots, X_k$.
- While we require that the response variable Y will be continuous, or approximately continuous. The *predictor variables* $X_1, X_2, \ldots, X_k$ can be either continuous or discrete.
- It is fairly standard notation to reserve k for the number of predictor variables in the regression model, and p for the number of parameters (regression coefficients or βs).
- When formulating a regression model, we want to explain the variation in the response variable by the variation in the predictor variable.

Statistical Assumptions for OLS Regression

There are two primary assumptions for OLS regression:

1. The regression model can be expressed in the form

   $$Y = \beta_0 + \beta_1 X_1 + \ldots + \beta_k X_k + \epsilon$$

   Notice that the model formulation specifies error term ϵ to be additive, and that the model parameters (βs) enter the modeling linearly, that is, β_i represents the change in Y for a one unit increase in X_i when X_i is a continuous predictor variable. Any statistical model in which the parameters enter the model linearly is referred to as a *linear model*.

2. The response variable Y is assumed to come from an independent and identically distributed (iid) random sample from a $N(\mathbf{X}\beta, \sigma^2)$ distribution where the variance of σ^2 is a fixed but unkown quantity. The statistical notation for this assumption is $Y \ N(\mathbf{X}\beta, \sigma^2)$.

Linear Versus Nonlinear Regression

Remember that a *linear model* is linear in the parameters, not the predictor variables.

- The following regression models are all linear regression models:

  $$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \epsilon$$

  $$Y = \beta_0 \beta_1 \ln(X_1) + \epsilon$$

- The following regression models are all nonlinear regression models:

  $$Y = \beta_0 \exp(\beta_1 X_1) + \epsilon$$

  $$Y = \beta_0 + \beta_2 \sin(\beta_1 X_1) + \epsilon$$

- If you know a little calculus, then there is an easy mathematical definition of a nonlinear regression model. In a nonlinear regression model at least one of the partial derivatives will be dependent on a model parameter.
- Any quantity that has a β in front of it counts as a degree of freedom used, and subsequently counts as a predictor variable.
- A hint to identify a nonlinear model is when a parameter is within a function, specifically a nonlinear function.

Distributional Assumptions for OLS Regression

The assumption $Y\ N(\mathbf{X}\beta, \sigma^2)$ can also be presented in terms of the error term ϵ. Most introductory bookos present the distributional assumption in terms of the error term ϵ, but more advanced books will use the standard Generalized Linear Model (GLM) presentation in terms of the response variable Y.

In terms of the error term ϵ the distributional assumption can also be presented as:

- The error term $\epsilon\ N(0, \sigma^2)$. Since $Y\ N(\mathbf{X}\beta, \sigma^2)$, then $\epsilon = Y - \mathbf{X}\beta$ has a $N(0, \sigma^2)$.

Distributional Assumptions in Terms of the Error

1. The errors are normally distributed.
2. The errors are mean zero.
3. The errors are independent and identically distributed (iid).
4. The errors are *homoscedastic*, i.e. the errors do not have any correlation "in time or space".

When we build statistical models, we will check the assumptions about the errors by assessing the model *residuals*, which are our estimates of the error term.

Homoscedasticity: a sequence or vector of random variables is *homoscedastic* if all random variables in the sequence or vector have the same finite variance. This is also known as the *homogeneity of variance*. [Wikipedia 2015a]

You'd need a pretty gross violation of *homoscedastic* in the kind of problems that we work with today.

Further Notation and Details

When we estimate an OLS regression model, we will be working with a random sample of response variables $Y_1, Y_2, \ldots, Y_n$, each with a vector of predictor variables $[X_{1i}, X_{2i}, \ldots, X_{ki}]$. In matrix notation we will denote the regression problem by

$$Y_{(n\times 1)} - X_{(n\times p)}\beta_{(p\times 1)} + \epsilon_{(n\times 1)}$$

where the matrix size is denoted by the subscript. Note that $X = [1, X_1, X_2, \ldots, X_k]$ and $\beta = [\beta_0, \beta_1, \beta_2, \ldots, \beta_k]$.

- When we want to express the regression in terms of a single observation, the new typically use the i subscript notation

  $$Y_i = \mathbf{X_i}\beta + \epsilon_i$$

  or simply

  $$Y_i = \beta_0 + \beta_1 X_{1i} + \ldots + \beta_k X_{ki} + \epsilon_i$$

Estimation and Inference for Ordinary Least Squares Regression

Notes below are from the following sources; [Bhatti 2011c].

It's important to understand some aspects of estimation and inference for every statistical method that is used.

Estimation - Simple Linear Regression

- A *simple linear regression* is the special case of an OLS regression model with a single predictor variable.

$$Y = \beta_0 + \beta 1 X + \epsilon$$

- For the ith observation we will denote the regression model by

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

- For the random sample $Y_1, Y_2, \ldots, Y_n$ we can estimate the parameters β_0 and β_1 by minimizing the sum of the squared errors,

$$\min \sum_{i=1}^{n} \epsilon_i^2$$

  which is equivalent to minimizing

$$\min \sum_{i=1}^{n} (Y_i - \beta_0 - \beta_1 X_i)^2$$

Estimators and Estimates for Simple Linear Regression

- The estimators for β_0 and β_1 can be computed analytically and are given by

$$\hat{\beta_1} = \frac{\sum (Y_i - \bar{Y})(X_i - \bar{X})}{(X_i - \bar{X})^2} = \frac{\text{Cov}(Y, X)}{\text{Var}(X)}$$

  and

$$\hat{\beta_0} = \bar{Y} - \hat{\beta_1} \bar{X}$$

- The regression line always goes through the centroid $(\bar{X}, \bar{Y})$.
- We refer to the formulas for $\hat{\beta_0}$ and $\hat{\beta_1}$ as estimators and the values that these formulas can take for a given random sample as the estimates.
- In statistics we put hats on all estimators and estimates.
- Given $\hat{\beta_0}$ and $\hat{\beta_1}$ the predicted value or fitted value is given by

$$\hat{Y} = \hat{\beta_0} + \hat{\beta_1} X$$

Estimation - The General Case

- We seldom build regression models with a single predictor variable. Typically we have multiple predictor variables denoted by $X_1, X_2, \ldots, X_k$, and hence the standard regression case is sometimes referred to as *multiple regression* in introductory regression texts.
- We can still think about the estimation of $\beta_0, \beta_1, \beta_2, \ldots, \beta_k$ in the same manner as the sime linear regression case

$$\min \sum_{i=1}^{n} (Y_i - \beta_0 - \beta_1 X_{1i} - \beta_2 X_{2i} - \ldots - \beta_k X_{ki})^2$$

  but the computations will be performed as matrix computations.

General Estimation - Matrix Notation

Before we set up the matrix formulation for the OLS model, let's begin by defining some matrix notation.

- The error vector $\epsilon = [\epsilon_1, \ldots, \epsilon_n]^T$.
- The response vector $Y = [Y_1, \ldots, Y_n]^T$.
- The design matrix or predictor matrix $X = [1, X_1, X_2, \ldots, X_k]$.
- The parameter vector $\beta = [\beta_0, \beta_1, \beta_2, \ldots, \beta_k]^T$.
- All vectors are column vectors, and the superscript T denotes the vector or matrix *transpose*.

General Estimation - Matrix Computations

- We minimize the sum of the squared error by minimizing $S(\beta) = \epsilon^T \epsilon$ which can be re-expressed as
$$S(\beta) = (Y - X\beta)^T (Y - X\beta)$$
- Taking the matrix derivative of $S(\beta)$, we get
$$S_\beta(\hat{\beta}) = -2X^T Y + 2X^T X\hat{\beta}$$
- Setting the matrix derivative to zero, we can write the expression for the least squares *normal equations*
$$X^T X\hat{\beta} = X^T Y$$
, which yield the estimator
$$\hat{\beta} = (X^T X)^{-1} X^T Y$$
- The estimator form $\hat{\beta} = (X^T X)^{-1} X^T Y$ assumes that the inverse matrix $(X^T X)^{-1}$ exists and can be computed. In practice your statistical software will directly solve the normal equations using a QR Factorization.

normal equations: projection of a linear space into a subspace to ensure that a solution exists.

QR Factorization: or QR decompositionof a matrix is a decomposition of a matrix A into a product $A = QR$ of an orthogonal matrix Q and an upper triangular matrix R [Wikipedia 2015b].

Statistical Inference with the t-Test

- In OLS regression the statistical inference for the individual regression coefficients can be performed using a t-test.

t-test: any statistical test using a t-statistic to derive the test and the p-value for the test. Alternatively, any statistical test that uses a t-statistic as the decision variable.

statistical test: have a null and alternative hypothesis, and a test statistic with a known distribution.

- When performing a t-test there are three primary components: (1) stating the null and alternative hypotheses, (2) Computing the value of the test statistic, and (3) deriving a statistical conclusion based on a desired significance level.
- Step 1: The null and alternate hypotheses for β_i are given by
$$H_0 : \beta_i = 0 \text{ versus } H_1 : \beta_i \neq 0$$
- Step 2: The t statistic for β_i is computed by
$$t_i = \frac{\hat{\beta}_i}{SE(\hat{\beta}_i)}$$
and has a degrees of freedom equal to the sample size minus the number of model parameters, i.e. $df = n - dim(model)$. For example if you had a regression model with two predictor variables and an intercept estimated on a sample of size 50, then the t statistic would have 47 degrees of freedom.
- Step 3: Reject the H_0 or Fail to Reject H_0 based on the value of your t statistic and your significance level. This decision can be made by using the p-value of your t statistic or by using the critical value for your significance level.

Confidence Intervals for Parameter Estimates

An alternative to performing a formal hypothesis test is to use a confidence interval for your parameter estimate. There is a duality between confidence intervals and formal hypothesis testing for regression parameters.

- The confidence interval for $\hat{\beta}_i$ is given by

$$\hat{\beta}_i \pm t(df, \frac{\alpha}{2}) \times SE(\hat{\beta}_i)$$

  where $t(df, \frac{\alpha}{2})$ is a t value from a theoretical t distribution, not a t statistic value.

- If the confidence interval does not contain zero, then this is the equivalent to rejecting the null hypothesis $H_0 : \beta_i = 0$.

Statistical Intervals for Predicted Values

The phrase *predicted value* is used in statistics to refer to the in-sample *fitted values* from the estimated model or to refer to the out-of-sample *forecasted values*. The dual use of this phrase can be confusing. A better habit is to use the phrase *in-sample fitted values* and in the *out-of-sample predicted values* to clearly reference these different values.

inference is an in-sample activity, measuring the quality of the model based on in-sample performance. *predictive modeling* is an out-of-sample activity, measuring the quality of the model based on out-of-sample.

- Given $\hat{\beta} = (X^T X)^{-1} X^T Y$ the vector of fitted values can be computed by $\hat{y} = X\hat{\beta} = HY$, where $H = X(X^T X)^{-1} X^T$. The matrix H is called the *hat matrix* since it puts the hat on Y.
- The point estimate $\hat{Y_0}$ at the point x_0 can be computed by $\hat{Y_0} = x_0^T \hat{\beta}$.
- The confidence interval for an in-sample point x_0 on the estimated regression function is given by

$$x_0^T \hat{\beta} \pm \hat{\sigma} \sqrt{x_0^T (X^T X)^{-1} x_0}$$

- The prediction interval for the point estimator $\hat{Y_0}$ for an out-of-sample x_0 is given by

$$x_0^T \hat{\beta} \pm \hat{\sigma} \sqrt{1 + x_0^T (X^T X)^{-1} x_0}$$

- Note that the out-of-sample prediction interval is always wider than the in-sample confidence interval.

Further Notation and Details

In order to compute the t statistic you need the standard error of the parameter estimate. Most statistical software packages should provide this estimate and compute this t statistic for you However, it is always a good idea to know from where this number comes. Here are the details needed to compute the standard error for $\hat{\beta}_i$.

- The estimated parameter vector $\hat{\beta}$ has the covariance matrix given by

$$\text{Cov}(\hat{\beta}) = \hat{\sigma}^2 X^T X$$

  where

$$\hat{\sigma}^2 = \frac{SSE}{n - k - 1}$$

- The variance of $\hat{\beta}_i$ is the ith diagonal element of the covariance matrix

$$\text{Var}(\hat{\beta}_i) = \hat{\sigma}^2 (X^T X)_{ii}$$

Analysis of Variance and Related Topics for Ordinary Least Squares Regression

Notes below are from the following sources; [Bhatti 2011d].

The ANOVA Table for OLS Regression

The Analysis of Variance or ANOVA Table is a fundamental output from a fitted OLS regression model. The output from the ANOVA table is sued for a number of purposes:

- Show the decomposition of the total variation
- Compute the R-Squared and Adjusted R-Squared metrics
- Perform the Overall F-test for a regression effect
- Perform a F-test for nested models as commonly used in forward, back-ward, and stepwise variable selection

Decomposing the Sample Variation

- The Total Sum of Squares is the total variation in the sample
- The Regression Sum of Squares is the variation in the sample that has been explained by the regression model
- The Error Sum of Squares is the variation in the sample that cannot be explained

SST	$\sum_i^n (Y_i - \bar{Y})^2$	Total Sum of Squares
SSR	$\sum_i^n (\hat{Y}_i - \bar{Y})^2$	Regression Sum of Squares
SSE	$\sum_i^n (Y_i - \hat{Y})^2$	Error Sum of Squares

Metrics for Goodness-Of-Fit in OLS Regression

The Coefficient of Determination - R-Squared

$$R^2 = \frac{SSR}{SST} = 1 - \frac{SSE}{SST}$$

- The Coefficient of Determintation R^2 will take values $0 \leq R^2 \leq 1$ and represents the proportion of the variance explained by the regression model.
- Implicitly, R^2 is a function of the number of parameters in the model. For a nested subset of predictor variables $p_0 < p_1$, i.e. p_1 contains the original p_0 predictor variables and some new predictor variables, R^2 will have a monotonic relationship such athat $R^2(p_0) \leq R^2(p_1)$.

Adjusted R-Squared

$$R^2_{ADJ} = 1 - \frac{\frac{SSE}{(n-k-1)}}{\frac{SST}{(n-1)}} = 1 - \frac{\frac{SSE}{(n-p)}}{\frac{SST}{n-1}}$$

- Note that the standard regression notation uses k for the number of predictor variables included in the regression model and p for the total number of parameters in the model. When the model includes an intercept term, then $p = k + 1$. When the model does not include an intercept term, then $p = k$.
- The Adjusted R-Squared metric accounts for the model complexity of the regression model allowing for models of different sizes to be compared.
- The Adjusted R-Squared metric will not be monotonic in the number of model parameters.
- The Adjusted R-Squared metric will increase until you reach an optimal model, then it will flatten out and likely decrease.

The Overall F-Test for a Regression Effect

Consider the regression model

$$Y = \beta_0 + \beta_1 X_1 + \ldots + \beta_k X_k$$

The Overall F-Test for a regression effect is a joint hypothesis test that at least one of the predictor variables has a non-zero coefficient.

- The null and alternate hypotheses are given by

  $$H_0 : \beta_1 = \ldots = \beta_k = 0 \text{ versus } H_1 : \beta_i \neq 0$$

  for some $i \in 1, \ldots, k$.

- The test statistic for the Overall F-test is given by

  $$F_0 = \frac{\frac{SSR}{k}}{\frac{SSE}{(n-p)}}$$

  which has a F-distribution with $(k, n-p)$ degrees-of-freedom for a regression model with k predictor variables and p total parameters. When the regression model includes and intercept, then $p = k + 1$. If the regression model does not include an intercept, then $p = k$.

- In some cases this can be very useful, such as if we had a categorical variable that has segmentation, the F-test can be useful. It is less likely that continuous variables will all have a zero coefficient.

The F-Test for Nested Models

For our discussion of nested models, let's consider two concrete examples which we will refer to as the *full model* (FM)

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3$$

and a *reduced model* (RM)

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2$$

Notice that the predictor variables in the reduced model are a subset of the predictor variables in the full model, i.e. $RM \subset FM$.

- In this notation we say that the FM *nests* the RM, or the RM is *nested by* the FM.
- We only use the terms *full model* and *reduced model* in the context of nested models.
- We can use a F-test for nested models to decide whether or not to include an additional predictor variable in the final model.

Given a *full model* and a *reduced model* we can perform a F-test for nested models for the exclusion of a single predictor variable or multiple predictor variables.

In the context of our example, we could test either of these null hypotheses:

- Example 1: Test a Single Predictor Variable

  $$H_0 : \beta_3 = 0 \text{ versus } H_1 : \beta_3 \neq 0$$

- Example 2: Test Multiple Predictor Variables

  $$H_0 : \beta_2 = \beta_3 = 0 \text{ versus } H_1 : \beta_i \neq 0$$

  for some $i \in 2, 3$.

The test statistic for the F-test for nested models will always have this form in terms of the FM and RM.

- Test Statistic for the Nested F-Test
$$F_0 = \frac{\frac{[SSE(RM) - SSE(FM)]}{(dim(FM) - dim(RM))}}{\frac{SSE(FM)}{[n - dim(FM)]}}$$
- The test statistic is based on the reduction in the SSE obtained from adding additional predictor variables. Note that $SSE(FM)$ is always less than $SSE(RM)$.
- The *dimension* of a statistical model is the number of parameters.

Connection to Forward Variable Selection

The F-test for nested models is a the standard statistical test implemented in most statistical software packages for performing forward and backward, and hence stepwise, variable selection.

Forward Variable Selection

- Given the model $Y = \beta_0 + \beta_1 X_1$ and a set of candidate predictor variables $Z_1, \ldots, Z_s$, how do we select the best Z_i to include in our model as X_2?
- In forward variable selection the FM will be $Y = \beta_0 + \beta_1 X_1 + \beta_2 Z_i$ and the RM will be $Y = \beta_0 + \beta_1 X_1$. The forward variable selection algorithm will select the Z_i with the largest F-statistic that is statistically significant at a predetermined level. The algotihm will continue to add predictor variables until there are no predictor variables that are statistically significant to the predetermined level.

Connection to Backward Variable Selection

Backward Variable Selection

- Given the model
$$Y = \beta_0 + \beta_1 X_1 + \ldots + \beta_s X_s$$
how do we eliminate the predictor variables whose effects are not statistically significant?
- In backward variable selection the FM will be $Y = \beta_0 + \beta_1 X_1 + \ldots + \beta_s X_s$ and the RM will be $Y = \beta_0 + \beta_1 X_1 + \ldots + \beta_{s-1} X_{s-1}$, for notional convenience. The backward variable selection algorithm will drop the X_i with the smallest F-statistic that is not statistically significant at a predetermined level. The algorithm will continue to drop predictor variables until there are no predictor variables that aren't statistically significant to the predetermined level.
- Note that both the forward and backward variable selection procedures consider only one variable at each iteration.

Statistical Inference Versus Predictive Modeling in OLS Regression

Notes below are from the following sources; [Bhatti 2011e]

- There are two reasons to build statistical models: (1) for inference, and (2) for prediction.
- Statistical inference is focused on a set of formal hypotheses, denoted by H_0 for the *null hypothesis* and H_1 for the *alternate hypothesis*, and a test statistic with a known sampling distribution. A test statistic will have a specified distribution, e.g. the t-statistic for an OLS regression parameter has a t-distribution with the degrees-of-freedom equal to $n - p$ where p is the number of model parameters for the dimension of the model.
- Predictive modeling is focused on accurately producing an estimated value for the primary quantity of interest or assigning an observation to the correct class (group). Typically, when we use the term 'predictive', we are referring to the model's ability to predict future or out-of-sample values, not in-sample values.

The Standard Modeling Process

0. Data Quality Check
1. Exploratory Data Analysis: How do our predictor variables relate to the response variable?
2. Model Identification: Which predictor variables should be included in our model?
3. Model Validation: Should we trust our models and the conclusions that we wish to derive from our model?

How we perform the Model Validation step is determined on the prescribed use of the model. Is the model to be used for statistical inference or is it to be used for predictive modeling?

Model Validation for Statistical Inference

- Model validation when the model is to be used for statistical inference is generally referred to as the *assessment of goodness-of-fit*.
- When we fit a statistical model, we have underlying assumptions about the probabilistic structures for that model. All of our statistical inference is derived from those probabilistic assumptions. Hence, if our estimated model, which is dependent upon the sample data, does not conform to these probabilistic assumptions, then our inference will be incorrect.
- When we validate a statistical model to be used for statistical inference, we are validating that the estimated model conforms to these probabilistic assumptions.
- For example in OLS regression we examine the residuals to make sure that they have a normal probability distribution and that they are homoscedastic.

Model Validation for Predictive Modeling

- Model validation when the model is to be used for predictive modeling is generally referred to as the *assessment of predictive accuracy*.
- When we fit a statistical model for predictive modeling, we can be much more tolerant of violations of the underlying probabilistic assumptions.
- Our primary interest in predictive modeling is estimating the response variable Y as 'accurately' as possible. When validating a predictive model, we tend to focus on summary statistics based on the quantity $(Y_i - \hat{Y}_i)$. Examples include the Mean Absolute Error (MAE) and the Mean Squared Error (MSE).
- The evaluation of predictive models is typically performed through a form of *cross-validation* where the sample is split into a *training sample* and a *test sample*. In this mdel validation, the model is estimated on the *training sample* and then evaluated out-of-sample on the *testing sample*.

Goodness-Of-Fit Versus Predictive Accuracy

- Goodness-Of-Fit
  - Goodness-Of-Fit (GOF) is assessed in-sample
  - The objective is to confirm the model assumptions
  - In OLS regression the GOF is typically assessed using graphical procedures (scatterplots) for the model residuals $e_i = Y_i - \hat{Y}_i$.
- Predictive Accuracy
  - Predictive Accuract (PA) is assessed out-of-sample
  - The objective is to measure the error of the predicted values
  - In OLS regression PA is typically assessed using error based metrics: Mean Square Error, Root Mean Square Error, and Mean Absolute Error.

Assessing the Goodness-Of-Fit in OLS Regression

- Validate the normality assumption: produce a Quantile-Quantile plot (QQ-Plot) of the residuals to compare their distribution to a nromal distribution.
- Validate the homoscedasticity assumption (equal variance): produce a scatterplot of the residuals against each predictor variable. If there is any structure in this plot, then the model will need a transformation of the predictor variable or an additional predictor variable added to the model.
- Interpret the R-Squared measure for your model. Applications tend to have typical ranges for "good" R-Squared values. If Model 1 has R-Squared of 0.23 and Model 2 has R-Squared of 0.54, then Model 2 should be preferred to Model 1, provided that Model 2 satisfies the other GOF conditions.
- By itself R-Square is not an exclusive measure of GOF. It's a measure of GOF provided everything else is satisfied.

Statistical Inference in OLS Regression

If our Analysis of Goodness-Of-Fit for our OLS regression does not uncover any major violations of the underlying probabilistic assumptions, then we can feel confident in our use of the two primary forms of statistical inference in OLS regression.

- The t-test for the individual model coefficients:

  $$H_0 : \beta_i = 0 \text{ versus } H_1 : \beta_i \neq 0$$

  for model coefficient i.
- The test statistic for the corresponding t-test is given by

  $$t_i = \frac{\hat{\beta}_i}{SE(\hat{\beta}_i)}$$

  where t_i has degrees of freedom equal to the sample size minus the number of model parameters, i.e. $df = n - dim(Model)$.

In addition to the 'local' tests of a regression effect for the individual predictor variables, we also have a 'global' test for a regression effect.

- The Overall F-test for a regression effect:

  $$H_0 : \beta_1 = \beta_2 = \ldots = 0 \text{ versus } H_1 : \beta_i \neq 0$$

  for some i, i.e. at least one of the predictor variables has an estimated coefficient that is statistically different from zero.

- The test statistic for the Overall F-test is given by:

$$F_0 = \frac{\frac{SSR}{k}}{\frac{SSE}{(n-p)}}$$

  which has a F-distribution with $(k, n-p)$ degrees-of-freedom for a regression model with k predictor variables and p total parameters. When the regression model includes an intercept, then $p = k+1$. If the regression model does not include an intercept, then $p = k$.

Predictive Accuracy in OLS Regression

The two primary metrics for assessing statistical models for out-of-sample predictive accuracy are Mean Square Error and Mean Absolute Error.

- Mean Square Error (MSE)

$$MSE = \frac{1}{n}\sum_{i=1}^{n}(Y_i - \hat{Y}_i)^2$$

  Root Mean Square Error (RMSE) is the square root of the MSE. There is no statistical reason to prefer one measure over the other. However, the RMSE can be used for presentation purposes when the MSE is very small or very large as the square root transformation will increase the small numbers and decrease the large numbers.

- Mean Absolute Error (MAE)

$$MAE = \frac{1}{n}\sum_{i=1}^{n}|Y_i - \hat{Y}_i|$$

The Bias-Variance Trade-Off

An interesting and useful property of Mean Square Error (MSE) is that it can be decomposed into two components: the prediction variance and the square of the prediction bias. This decomposition is referred to as the *Bias-Variance Trade-Off*, and it is referenced throughout predictive modeling, especially in the presentation of concepts from statistical and machine learning.

- Throughout these notes we have been using the *empirical* Mean Square Error for the predictved values $\hat{Y}_i$.

$$MSE = \frac{1}{n}\sum_{i=1}^{n}(Y_i - \hat{Y}_i)^2$$

- The *Bias-Variance Trade-Off* is presented from the *theoretical Mean Square Error*

$$MSE = \mathbb{E}(Y_i - \hat{Y}_i)^2$$

  where $\mathbb{E}[X]$ denotes the mathematical expectation of X.

Final Comments on the Bias-Variance Trade-Off

The crux of the *Bias-Variance* Decomposition is to note that both terms of the decomposition are non-negative. Hence, we can choose to minimize either the Variance or the bias.

- The variance of the predicted value is a measure of the spread of the predicted value from its mean.
- The bias of the predicted value is a measure of the distance from the mean of the predicted value to the target value.

Both of these components are functions of *model complexity*, i.e. the nubmer of parameters in the model. Ideally, you would want to have your prediction to be accurate (low bias) and precise (low variance). Bias will decline and variance will increase as the model complexity increases.

Further Notation and Details

The Mean Square Error of the predicted values $\hat{Y}_i MSE = \frac{1}{n}\sum_{i=1}^{n}(Y_i - \hat{Y}_i)^2\$$ should not be confused with the estimate or variance parameter σ^2 in an OLS regression model with the Error Sum of Squares denoted by SEE and p parameters,

$$\sigma^2 = \frac{SSE}{n-p}$$

which is frequently referred to as the *mean square error* of the regression or the *mean square* of the residuals, but is not denoted by MSR as to not be confused with the *mean square of the regression* ($MSR = \frac{SSR}{k}$).

If you are in the context of a fitted OLS regression model, then the term MSE is referring to the estimate $\hat{\sigma^2}$.

Introduction to Principle Components Analysis

Notes below are from the following sources; [Bhatti 2011f]

- *Statistical Interpretation* - PCA is a transformation of a set of correlated random variables to a set of uncorrelated (or orthogonal) random variables
- *Linear Algebra Interpretation* - PCA is a rotation of the coordinate system to the canonical coordinate system, i.e. the natural coordinate system defined by the variation in the data.
- *Numerical Linear Algebra Interpretation* - PCA is a reduced rank matrix approximation that facilitates dimension reduction.

Facts and Caveats

- PCA does not require any statistical assumptions, e.g. the data are not assumed to have a multivariate normal distribution.
- PCA is a (numerical) linear algebra technique, i.e. it relies on a matrix factorization (the Spectral Decomposition or Singular Value Decomposition).
- PCA is sensitive to the scale of the data. Most of the time the data should be standardized, i.e. the variables should have a (0,1) distribution. When the data re standardized our covariance matrix and correlation matrix are the same matrix.
- If the data are 'standardized' to a common scale that is not (0,1), then it should not be standardized to a (0,1) distribution.

Why do we use PCA?

- PCA can be used in its own right to understand the covariance structure in multivariate data with respect to the measured basis.
- PCA can be used as a method to create a reduced rank approximation to the covariance structure, i.e. PCA can be used to approximate the variation in p predictor variables using $k < p$ principal components. This property is typically referred to as dimension reduction.
- PCA can be used as a means of creating a set of orthogonal predictor variables from a set of raw predictor variables. Since the principal components created from the original predictor variables are orthogonal, we can use PCA as a remedy for multicollinearity in regression problems or as a preconditioner to cluster analysis.

How do we compute the principal components?

- Consider the $n \times p$ data matrix of predictor variables $X = [X_1, \ldots, X_p]$.
- Depending on your software the data may need to be standardized before the principal components are computed. This is typically true if you use a software to compute eigenvalues and eigenvectors. Statistical software designed to perform PCA, such as princomp in SAS, will typically internally standardize the data for you.
- Compute the eigenvalue-eigenvector pairs $(\lambda_1, e_1), \ldots, (\lambda_p, e_p)$ of the square matrix $X^T X$ where the eigenvalues are ordered largest to smallest such that $\lambda_i > \lambda_j$ for $i > j$.
- Your software will compute the eigenvalue-eigenvector pairs using a matrix factorization called Singular Value Decomposition or SVD.
- Compute the principal components $Z_1, \ldots, Z_p$ using the eigenvalues as the component loadings
- In vector format we can compute each component individual

$$Z_i = X \times e_i$$

or we can compute all of the principal components using one matrix computation

$$[Z_1, \ldots, Z_p] = X \times [e_1, \ldots, e_p]$$

.

How many principal components should we use?

- A $p \times p$ matrix will yield p principal components if all of the eigenvalues are non-zero.
- One standard approach to selecting the number of principal components to keep is to use the *scree* plot. The *scree* plot plots the number of components on the x-axis against the proportion of the variance explained on the y-axis. The suggested number of principal components to keep is the number where the plot forms an 'elbow', i.e. the point where the curve starts to flatten out.
- Another rule for selecting the number of principal components to keep is to use the minimum eigenvalue rule. A frequently used rule is the *Kaiser Rule*, which recommends that the number of principal components to keep is equal to the number of eigenvalues grater than one.
- Other rules exist and ad hoc decisions can be made. Keep in mind that in some problems you might keep all of the principal components.
  - Example: Keep at least as many principal components needed to explain at least 70% of the total variation in the data.

How do I know if I have kept the correct number of principal components?

- Frequently the scree plot will present some ambiguity in the number of components to keep, e.g. should I keep four or five principle components?
- The 'correct' number of principal components to keep will depend on the application. If you are using PCA as a preconditioner for regression analysis or cluster analysis, then the effectiveness of these applications under the alternate choices would determine which number is the best to keep. In this sense the unsupervised learning problem has been transformed into a supervised learning problem.
- If the PCA is not directly tied to any application, then the choice of the number of components to keep is always heuristic. Formal inference for the number of components is available under a multivariate normal distribution assumption.

Exploratory Factor Analysis

Notes below are from the following sources; [Bhatti 2011g]

- Factor analysis is a statistical modeling technique used to model the covariance structure in multivariate data.
- Factor analysis is a statistical modeling technique for estimating unobserved (or latent) relationships using observed (or measured) variables.
- As a statistical modeling technique, factor analysis has statistical assumptions

Why do we use Factor Analysis

- We use FA to model the correlation structure in a set of measured variables.
  - If your data is properly scaled the covariance matrix and correlation matrix are the same matrix.
- We use FA to facilitate a dimension reduction from the observed measurement variables to the smaller set of unobserved latent factors.
- We use FA to improve the interpretability of our multivariate data.

When should we use Factor Analysis

Factor analysis is most useful on problems that are 'natural factor analysis problems'. It would be rare to use factor analysis on a bunch of random data that was just pulled without obtaining understanding of the data space. In that matter Factor Analysis is very different than Principal Components Analysis. You can apply PCA to any data that is continuous or approximately continuous.

- All variable names are known, i.e. we know and understand all of the measurement variables
- All measurement variables have been purposely selected under the guidance that they represent a measurement of a quality or trait that is recognized as important but that cannot be directly measured.
- We are able to obtain multiple measurements for each of the unmeasurable qualities (the latent factors).

Examples of Natorual Factor Analysis Problems

- Measurement of physical attributes: speed, strength, agility
- Measurement of educational attainment: math, reading, problem solving
- Measurement of personality: rational, social, empathetic

The discipline of psychology uses factor analysis a lot. You can find many examples of factor analysis in psychometric literature.

Exploratory versus Confirmatory Factor Analysis

- EFA is performed when we have no preconceived notions about the factor structure (i.e. the factor loadings) that may exist in a set of multivariate data. Since we have no preconceived notions of the factor structure, we are performing an exploratory data analysis, hence the name exploratory factor analysis.
- CFA is performed when we want to statistically test a specific factor structure. CFA will require the formal statistical assumptions of maximum likelihood estimation so that formal statistical inference can be applied to the data to draw statistical conclusions. CFA is related to other topics such as path analysis and simultaneous equations.

How does Factor Analysis Differ from Principal Components Analysis?

- FA is a statistical model while PCA is not.
- FA is a statistical model for the correlation structure in a multivariate data. PCA is not a model for the correlation structure. PCA is a rotation of the coordinate axes.
- FA is focused on producing a representation of the correlation structure that will provide an enhanced data interpretation.
- PCA can be used with anonymized variables, ie.e. data with no known names, while FA requires well defined and known variables. FA is more subjective than PCA, and also more focused on interpretability than PCA, and hence FA is difficult to use with a larger number of variables or anonymized variables.

A Modeling Process for Factor Analysis

- Since we are focused on statistical modeling, and not statistical theory, we will approach factor analysis from the modeling perspective.
- How do we define a modeling process for factor analysis?
- How different or similar will our modeling process for factor analysis be to our standard linear models approach?

A Strategy for Performing a Factor Analysis

- Perform a Principal Factor Analysis with a Varimax rotation.
- Perform an iterative Principal Fator Analysis with a Varimax rotation.
- Performa Maximum Liklihood Factor Analysis with a Varimax rotation.
- Compare the solutions from these three factor analyses:
  - Did each factor analysis yield roughly the same factor loadings?
  - Is one set of factors more interpretable than the other set?
- Evaluate the factor loadings over a range of *common factors*, i.e. instead of just looking at the results for $k = 4$ also consider values for k in $2, 3, 4, 5, 6$.
- If you have enough data, then evaluate your prospective factor loadings through bootstrapping or cross validation. As with any statistical relationship, for that relationship to represent a 'universal truth', then it must exist in many samples. We can effectively construct this ideal situation by employing either of these methods.

This strategy has been adapted from the advice provided in [Johnson et al. 1992].

Common Factor Model: General Assumptions

- Let $X_1, X_2, \ldots, X + p$ be observable random variables. In the context of FA we will call these variables the response variables.
  - These are the variables for which we will try to model the correlation structure.
- Let f_i denote an unobservable concept called a common factor.
- Let λ_{ij} denote the factor loading for the jth common factor f_j on the ith response X_i. Note that the factor loading λ_{ij} represents the correlation between the common factor f_j and the response variable X_i.
- Let u_j be the *specific error* or *unique factor*.
- In factor analysis we assume that each response X_i can be deconstructed into a set of *common factors* $f_1, f_2, \ldots, f_k$ and a *unique factor* u_i.

Thurstone's Common Factor Model

The Common Factor Model is defined by the following relationships:

$$X_1 = \lambda 11 f_1 + \lambda 12 f_2 +, \ldots + \lambda 1k f_k + u_1$$
$$X_1 = \lambda 21 f_1 + \lambda 22 f_2 +, \ldots + \lambda 2k f_k + u_2$$
$$X_1 = \lambda p1 f_1 + \lambda p2 f_2 +, \ldots + \lambda pk f_k + u_p$$

- Each response variable X_i is assumed to be related to the unobserved factors f_j through a linear equation with the coefficients (factor loadings) given by λ_{ij}. Since the relationship will not be exact, each linear equation has a response specific error called the *unique factor* or *specific error*.
- The factors $f_1, f_2, \ldots, f_k$ are called *common factors* because they are common to all of the response variables $X_1, X_2, \ldots, X_p$.
- The relationship between the response variables X_i and the *common factors* f_j looks very similar to an OLS regression model, except the common factors are not directly observable (measurable).
- Let's make sure we remember the objective of factor analysis. Specifically, we do not want to be able to reproduce the response variables X_i. When using factor analysis, we want to model the correlation structure of the response variables.
- How doe we define the correlation structure? We can describe the correlation structure of a multivariate data using either the correlation matrix or the covariance matrix.
- We can use the Common Factor Model defined above in matrix format:

$$\mathbf{X} = \mathbf{\Lambda f} + \mathbf{u}$$

- From this matrix representation of the Common Factor Model we can then define the covariance matrix with some additional statistical assumptions.
- Note: In theoretical discussions we discuss the correlation structure using the covariance matrix. However, in practice we use the correlation matrix. When the data are standardized these matrices are the same. When we perform data analysis, we (our software) will standardize the data in order to make these matrices the same.
- It might help to think of the matrix $\mathbf{\Sigma}$ as the covariance matrix for standardized data so that it can be interpreted as a correlation matrix.

Needed Additional Assumptions

To define the correlation structure we need some additional assumptions.

- The unique (specific) factor u_i has the unique (specific) variance ψ_i, i.e. $\text{Var}(u_i) = \psi_i$.
- The common factors f_j are mean zero, variance 1, and independent of each other, ie.e f_n and f_m are independent for all n not equal m.
- The unique (specific) factors and the common factors are independent of eachother, i.e. u_i and f_i are independent for all i and j.

Defining the Correlation Structure

Using the additional statistical assumptions we can define the correlation structure for the matrix $\mathbf{X}$.

$$\begin{aligned} \text{Cov}(\mathbf{X}) &= \text{Cov}(\mathbf{\Lambda} f + u) \\ &= \text{Cov}(\mathbf{\Lambda} f) + \text{Cov}(u) \\ &= \mathbf{\Lambda}\text{Cov}(f)\mathbf{\Lambda^T} + \text{Cov}(u) \end{aligned} \tag{1}$$

This relationship yields the commonly displayed correlation structure:

$$\mathbf{\Sigma} = \mathbf{\Lambda}\mathbf{\Lambda_T} + \mathbf{\Psi}$$

Where:

$$\mathbf{\Lambda} = \begin{bmatrix} \lambda_{11} & \dots & \lambda_{1k} \\ \vdots & \dots & \vdots \\ \lambda_{p1} & \dots & \lambda_{pk} \end{bmatrix}$$

and

$$\mathbf{\Psi} = \begin{bmatrix} \psi_1 & 0 & \dots & 0 & 0 \\ 0 & \psi_2 & 0 & \dots 0 & \\ \vdots & \dots & \dots & \dots & \vdots \\ 0 & 0 & 0 & 0 & \psi_p \end{bmatrix}$$

Correlation Components

The covariance matrix $\mathbf{\Sigma}$ has several components and related quantities of interest.

- The diagonal entries of the covariance matrix are individual variance terms.

$$\sigma_i^2 = \lambda_{i1}^2 + \ldots + \lambda_{ip}^2 + \psi_i$$

- The sum of the squared factor loadings is called the *common variance* or the communality and typically denotd by h_i^2.

$$h_i^2 = \lambda_{i1}^2 \ldots + \lambda_{ik}^2$$

- The term ψ_i is called the *unique* (specific) variance.
- The covariance between X_i and X_j is given by

$$\sigma_{ij} = \sum_{n=1}^{k} \lambda_{in}\lambda_{jn}$$

How do we estimate the Factor Loadings?

The correlation structure

$$\mathbf{\Sigma} = \mathbf{\Lambda}\mathbf{\Lambda_T} + \mathbf{\Psi}$$

has only one estimatable qunatity and two unknowns. How can we estimate the factor loadings $\mathbf{\Lambda}$?

We can estimate $\mathbf{\Sigma}$ from the data. If we had an estimate for $\mathbf{\Psi}$ then we could estimate the **reduced covariance matrix**

$$\hat{\mathbf{\Sigma}} - \hat{\mathbf{\Psi}} = \hat{\mathbf{\Lambda}}\hat{\mathbf{\Lambda}}^{\mathbf{T}}$$

How do we estimate $\mathbf{\Psi}$? Instead of estimating $\mathbf{\Psi}$ directly, most algorithms initiate the estimation with an estimate of the reduced covariance matrix $\hat{\mathbf{\Sigma}} - \hat{\mathbf{\Psi}}$, and then back out estimates for $\mathbf{\Psi}$.

We estimate the reduced covariance matrix $\hat{\mathbf{\Sigma}} - \hat{\mathbf{\Psi}}$ by computing the estimated covariance matrix $\hat{\mathbf{\Sigma}}$ and replacing the 1's on the diagonal with $\hat{h_1} - \hat{h_k}$ for a k factor model. These estimates $\hat{h_i}$ are called the **prior communality estimates**.

Estimating the Reduced Covariance Matrix

There are two common methods for estimating the reduced covariance matrix $\hat{\mathbf{\Sigma}} - \hat{\mathbf{\Psi}}$.

- The most common estimation procedure is to replace the diagonal of the estimated covariance matrix with the Square Multiple Correlation coefficients (the SMC priors option in SAS). Note that the Square Multiple Correlation coefficient is the R-Squared value from regressing one X_i on all of the other X_j.
- The other option is to replace the diagonal of the estimated covariance matrix with the absolute value of the maximum correlation coefficient for that row (The MAX priors option in SAS).

In either case the factor loadings are then computed by performing Principal Components Analysis (as a matrix factorization technique) on the reduced covariance matrix. This estimation technique is referred to as the *Principal Factor Analysis*.

Iterative Principal Factor Analysis

- Iterative Principal Factor Analysis uses an updated estimate for $\mathbf{\Psi}$ at each interaction.
- After any iteration we can compute
$$\hat{\Psi} = \hat{\mathbf{\Sigma}} - \hat{\mathbf{\Lambda}}\hat{\mathbf{\Lambda}}^{\mathbf{T}}$$
After updating the estimate $\mathbf{\Psi}$, the algorithm will recompute the factor loadings using the new reduced covariance matrix until the factor loadings 'stabilize'.
- When Principal Factor Analysis is computed using an iterative estimate for the $\mathbf{\Psi}$, the procedure is called *Iterative Principal Factor Analysis*.

Maximum Likelihood Factor Analysis

- Maximum Likelihood Factor Analysis is the statistical approach to factor analysis. The response varaibles X_i are assumed to have a multivariate normal distribution, and hence the covariance matrix $\mathbf{\Sigma}$ is assumed to have a Wishart distribution.
- Principal Factor Analysis and Iterative Principal Factor Analysis are not formal statistical models. They are not estimated from a likelihood function, and hence they do not have an means of formal inference.
- Maximum Likelihood Factor Analysis is the only formal estimation procedure for factor analysis, and hence the only estimation procedure with formal inference for factor loadings (confidence intervals) and statistical tests for goodness-of-fit.
- Inference includes a formal test of model adequacy for the number of factors and the use of AIC as a means of model selection for the number of factors.

Model Goodness-Of-Fit

What does it mean for a factor model to fit well?

- Unfortunately, in practice the goodness-of-fit of a factor model is typically completely determined by its interpretability.
- All communality estimates should be less than 1, i.e. no Heywood cases.
  - "Heywood cases" [are] negative estimates of variances or correlation estimates greater than one in absolute value. . . "
  - From publication [Heywood 1931]
- Factor loadings should exhibit a *simple factor structure*. This is ideal, but seldom hold in practice.
- Objective Cririques of Fit:
  - Small residual matrix: The estimated matrix $\hat{\mathbf{\Psi}}$ can be interpreted as a residual matrix. Componentwise metrics (matrix norms) such as Mean Absolute Error (MAE) and Mean Square Error (MSE) can be used to compare different factor models for a relative fit.
  - Statistical Inference: Use MLE and statistical inference to justify your factor model.

Caveats of Goodness-Of-Fit for Factor Analysis

- Without MLE all GOF is subjective, or at least relative.
- Using MLE requires the assumption of multivariate normality. This assumption is difficult to validate from the model. One would have to validate the assumption prior to using ML FA by looking at the marginal distributions.
- Software will frequently output estimates for the Heywood case, and some people will use this output. The validity of the output in the Heywood case will be contentious.
- One item that tends to be lost in the discussion of GOF is the overall legitimacy of your factor relationships. Ar ethey real, i.e. would they be apparent across many samples? Or are they sample specific, i.e. do they just happen to show in your sample? More importantly, did you deviate your sample from a random sample and generate a spurious relationship?

Factor Rotations

Factor rotations are a common practice in factor analysis. There are two types of factor rotations: (1) orthogonal rotations and (2) oblique rotations.

- **Orthogonal** rotations will yield orthogonal factors after the rotation. The most common orthogonal rotation is the **Varimax** rotation.
- **Oblique** rotations will yield correlated factors after the rotation. The most common oblique rotation is the **Promax** rotation.

Item to note:

- Factor rotations do not improve the 'fit' of the factor model, only the factor interpretation. Factor rotations are used as a means to obtain a 'simple structure', and hence be more interpretable.
- After applying an oblique rotation the loadings matrix no longer represents the correlation structure between the observed variables and the unobserved factors.

Rotations to Simple Structure

The primary reason to employ a factor rotation is to improve the interpretation of the factor structure by rotating to a **simple structure**.

- Each row of $\mathbf{\Lambda}$ should contain at least one zero.
- Each column of $\mathbf{\Lambda}$ should contain at least k zeros for a k factor model. Ideally, the factor grouping of the response variables is clearly identifiable.
- When comparing pairs of columns from $\mathbf{\Lambda}$, the pairs should have some elements that are zero in one column but nonzero in the other column. Again, ideally the factor groupings of the response variables are identifiable and unique to a group.

In practice simple factor structures are difficult to obtain. The user can try to fit simple factor structures through factor rotations and/or the inclusion or exclusion of response variables. In practice all of these decisions are subjective and considered to be the user's prerogative.

This advice been adapted from [Morrison 1990]

Caveats and Potential Shortcomings of Factor Analysis

- The common factor model is an *underdetermined* model (more unkowns than equations), and hence it has multiple solutions instead of a single unique solution. The existence of multiple solutions in any problem always causes confusion.
- Factor analysis will only work on a problem that is a 'factor analysis problem', i.e. one where the user have performed a set of measurements that are intended to be related, i.e. measure an underlying concept. In this sense FA is not a general method for dimension reduction.
- In practice your factor analysis results will never be as 'nice' as you think they should be.

Relationship Between the Covariance Matrix and the Correlation Matrix

Throughout the notes we have referenced the covariance matrix $\mathbf{\Sigma}$ and its estimate $\hat{\mathbf{\Sigma}}$ (denoted by $\mathbf{S}$ in some books). When your data are centered and scaled, then the covariance matrix and the correlation matrix are the same matrix. Hence, when your data are not standardized, the difference between the correlation matrix is a scaling matrix.

- For two random variables X_1 and X_2 with standard deviations σ_1 and σ_2 we have:

$$\text{Corr}(X_1, X_2) = \frac{\text{Cov}(X_1, X_2)}{\sigma_1 \sigma_2}$$

- Let $\mathbf{R}$ denote the correlation matrix. Consider the matrix representation of the problem for two random variables X_1 and X_2:

$$\mathbf{R} = \begin{bmatrix} \frac{\text{Cov}(X_1,X_1)}{\sigma_1^2} & \frac{\text{Cov}(X_1,X_2)}{\sigma_1\sigma_2} \\ \frac{\text{Cov}(X_2,X_1)}{\sigma_2\sigma_1} & \frac{\text{Cov}(X_2,X_2)}{\sigma_2^2} \end{bmatrix}$$

- In terms of $\mathbf{\Sigma}$ we can write $\mathbf{R}$ as:

$$\mathbf{R} = \text{diag}(\mathbf{\Sigma})^{-\frac{1}{2}} \times \mathbf{\Sigma} \times \text{diag}(\mathbf{\Sigma})^{-\frac{1}{2}}$$

- Since diagonal matrices are always invertible, any relationship specified with a covariance matrix could be respecified with a correlation matrix using an algebraic substitution. However, in practice we would simply perform factor analysis on the correlation matrix (or the software would do this automatically).
- Remember, although we presented all of the technical details in terms of the covariance matrix, in practice we are modeling the correlation matrix, and we are trying to find a factor structure to describe and reproduce the correlation matrix.

Study Questions for Ordinary Least Squares Regression

Question: When we refer to a 'simple linear regression', to what type of model are we referring? How does a 'simple linear regression' differ from a 'multiple regression'?

Response from [Montgomery et al. 2012] pages 2 and 4.

$$y = \beta_0 + \beta_1 x + \varepsilon$$

The equation abovce is called a **linear regression model**. Customarily x is called the independent variable and y is called the dependent variable. However, this often causes confustion with the concept of statistical independence, so we refer to x as the **predictor** or **regressor**and y as the **response** variable. Because the equation above involves only one **regressor** variable, it is called a **simple linear regression model**.

In general the response variable y may be related to k regressors, $x_1, x_2, \ldots, x_k$, so that:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_k x_x + \varepsilon$$

This is called a **multiple linear regression model** because more than one regressor is involved. The adjetive linear is employed to indicate that the model is linear in the parameters $\beta_0, \beta_1, \ldots, \beta_k$, not because the y is a linear function of the x's.

Question: In statistics, and in this course, we use the term 'regression' as a general term. What do we mean by the term 'regression'? What is the objective of a 'regression model'?

Response from [Montgomery et al. 2012] pages 1.

Regression analysis is a _statistical technique **for investigating and** modeling the relationship between variables___. The goal of regression analysis is to determine the values of parameters for a function that causes the function to best fit a set of data. Regression analysis helps one understand how the typical value of the dependent variable changes when any one of the independent variables is varied, while the other independent variables are held fixed [Wikipedia 2015c].

Question: What do we mean by 'linear regression'? represent a linear regression? Which equations represent a linear regression?

(a) $y = \beta_0 + \beta_1 x_1$
(b) $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2^2$
(c) $y = \beta_0 + \exp(\beta_1) x_1$

Both (a) and (c) are representations of linear regression because they only have a single **regressor**. (b), with the presence of x_2^2 has two **regressors**.

Question: Before building statistical models, it is a common and preferred practice to perform an Exploratory Data Analysis (EDA). What constitutes an EDA for a simple linear regression model? Is this EDA satisfactory for a multiple regression model, or do we need to change or extend the EDA? As we move forward in this course we will also learn about logistic regression models and survival regression models, will these methods need their own EDA or is EDA general to all statistical models?

Question: In the simple linear regression model what is the relationship between R-squared and the correlation coefficient rho?

Question: How do we interpret a regression coefficient in OLS regression?

Question: Frequently, as a form of EDA for OLS regression we make a scatterplot between the response variable Y and a predictor variable X. As an assumption of OLS, the response variable Y must be continuous. However, the predictor variable X could be continuous or discrete. When the predictor variable is discrete, does a scatterplot still make sense? If not, what type of visual EDA does make sense? Does the appropriateness of the scatterplot make sense if the discrete variable takes on many discrete values (such as the set of integers, think of dollar amounts rounded to the nearest dollar) versus only a few discrete values(such as a coded categorical variable which only takes the values 1, 2, or 3)?

Question: The simple linear regression model is a special case of 'Multiple Regression' or 'Ordinary Least Squares'(OLS) regression. (We will typically use the term OLS regression.) What are the assumptions of OLS regression? In the final step of a regression analysis we perform a 'check of model adequacy'. What model diagnostics do we use to validate our fitted model against the model assumptions of OLS regression?

Question: How are the parameters, i.e. the model coefficients, estimated in OLS regression? How does this relate to maximum likelihood estimation? How do you show the relationship between OLS regression and maximum likelihood estimation?

Question: What is the overall F-test? What is the null hypothesis and what is the alternate hypothesis? The overall F-test is also called the 'test for a regression effect'. Why is it called this?

Question: What is the difference between R-squared and adjusted R-squared? How is each measure computed, and which measure should we prefer? How does the interpretation of R-squared change as we move from the simple linear regression model to the multiple regression model?

Question: The simple linear regression model $Y = b_0 + b_1 * X_1$ has three parameters. Two of the parameters are b_0 and b_1. What is the third parameter?

Question: What is a sampling distribution? What theoretical distribution do the parameter estimates have in OLS regression? What distribution do we use in practice? Why do we use a different distribution in practice?

Question: The final step of a regression analysis is a 'check of model adequacy'. This 'check of model adequacy' or 'goodness-of-fit' is a very important step in regression analysis. Why? Which quantities in the regression output are affected when the fitted model deviates from the underlying assumptions of OLS regression?

Question: Nested Models: Given two regression models M1 and M2, what does it mean when we say that 'M2 nests M1''?

Question: What is the Analysis of Variance Table for a regression model? How do we interpret it and what statistical tests and quantities can be computed from it?

Question: When the intercept is excluded in a regression model, how does the computation and the interpretation of R-squared change? Fit a no intercept model in SAS and check the SAS output for any noted differences.

Question: How do we interpret the diagnostic plots output by the PLOTS(ONLY)=(DIAGNOSTICS) option in PROC REG in SAS?

Question: Why do we plot each predictor variable against the residual as a model diagnostic?

Question: Why do we perform transformations in the construction of regression models? Name at least two reasons.

Question: What is multicollinearity and how does it affect the parameter estimates in OLS regression? How do we diagnose multicollinearity?

Question: What is a Variance Inflation Factor (VIF) and how does it relate to multicollinearity?

Question: Given a set of predictor variables $X_1, \ldots, X_n$, which are determined to show a high degree of multicollinearity between some of the variables, how should we choose a subset of these predictor variables to reduce the degree of multicollinearity and improve our OLS regression performance?

Question: Variable Selection: How does forward variable selection work? How does backward variable selection work? How does stepwise variable selection work?

Study Questions for Multivariate Analysis

Principle Components Analysis

Question: Principal Components Analysis (PCA): What is principal components analysis? How does PCA eliminate the problem of multicollinearity? What does it mean for X1 and X2 to be orthogonal? In order to better understand orthogonality, take the building prices data set and perform these steps:

(a) Perform a PROC CORR on X1-X9.
(b) Create nine orthogonal predictor variables using PCA. Call these variables Z1-Z9.
(c) Perform a PROC CORR on Z1-Z9.

Question: Principal Components Analysis is described as a method of 'dimension reduction'. How does PCA reduce the dimension of a statistical problem? How do you select the reduced dimension for your problem.

Factor Analysis

Question: Are the factor scores always orthogonal? Are they orthogonal after a rotation?

Question: If two analysts perform a factor analysis, are they likely to arrive at the same result? If the same two analysts perform a principal components analysis, are they likely to get the same result?

Question: What is the first step in performing a factor analysis?

Question: In the context of factor analysis, what is the communality of factors?

Cluster Analysis

Question: What is the difference between hierarchical and non-hierarchical clustering?

Question: What is linkage? What types of linkage are there?

Question: How do we examine the goodness-of-fit of a cluster analysis or two comparative cluster analyses?

Question: Do the data need to be treated before we perform a cluster analysis?

References

Bhatti, D.C. 2011a. Statistical preliminaries and mathematical notation. http://nwuniversity.adobeconnect.com/p2u4z1zop3a/.

Bhatti, D.C. 2011b. Statistical assumptions for ordinary least squares regression. http://nwuniversity.adobeconnect.com/p14gl2rughs/.

Bhatti, D.C. 2011c. Estimation and inference for ordinary least squares regression. http://nwuniversity.adobeconnect.com/p9rbxcf6431/.

Bhatti, D.C. 2011d. Analysis of variance and related topics for ordinary least squares regression. http://nwuniversity.adobeconnect.com/p5hhosjmkbh/.

Bhatti, D.C. 2011e. Statistical inference versus predictive modeling in oLS regression. http://nwuniversity.adobeconnect.com/p5vo4o95h0r/.

Bhatti, D.C. 2011f. Introduction to principal components analysis. https://nwuniversity.adobeconnect.com/_a799312996/p32n9izw7y3.

Bhatti, D.C. 2011g. Exploratory factor analysis.

Heywood, H. 1931. On finite sequences of real numbers. *Proceedings of the Royal Society of London. Series A, Containing Papers of a Mathematical and Physical Character*, 486–501.

Johnson, R.A., Wichern, D.W., and others. 1992. *Applied multivariate statistical analysis*. Prentice hall Englewood Cliffs, NJ.

Montgomery, D.C., Peck, E.A., and Vining, G.G. 2012. *Introduction to linear regression analysis*. John Wiley & Sons.

Morrison, D.F. 1990. Multivariate statistical methods. 3. *New York, NY. Mc*.

Wikipedia. 2015a. Homoscedasticity — wikipedia, the free encyclopedia. /url\protect\T1\textbraceleftthttp://en.wikipedia.org/w/index.php?title=Homoscedasticity&oldid=650997386\protect\T1\textbraceright.

Wikipedia. 2015b. QR decomposition — wikipedia, the free encyclopedia. /url\protect\T1\textbraceleftthttp://en.wikipedia.org/w/index.php?title=QR_decomposition&oldid=655839188\protect\T1\textbraceright.

Wikipedia. 2015c. Regression analysis — wikipedia, the free encyclopedia. /url\protect\T1\textbraceleftthttp://en.wikipedia.org/w/index.php?title=Regression_analysis&oldid=647603059\protect\T1\textbraceright.